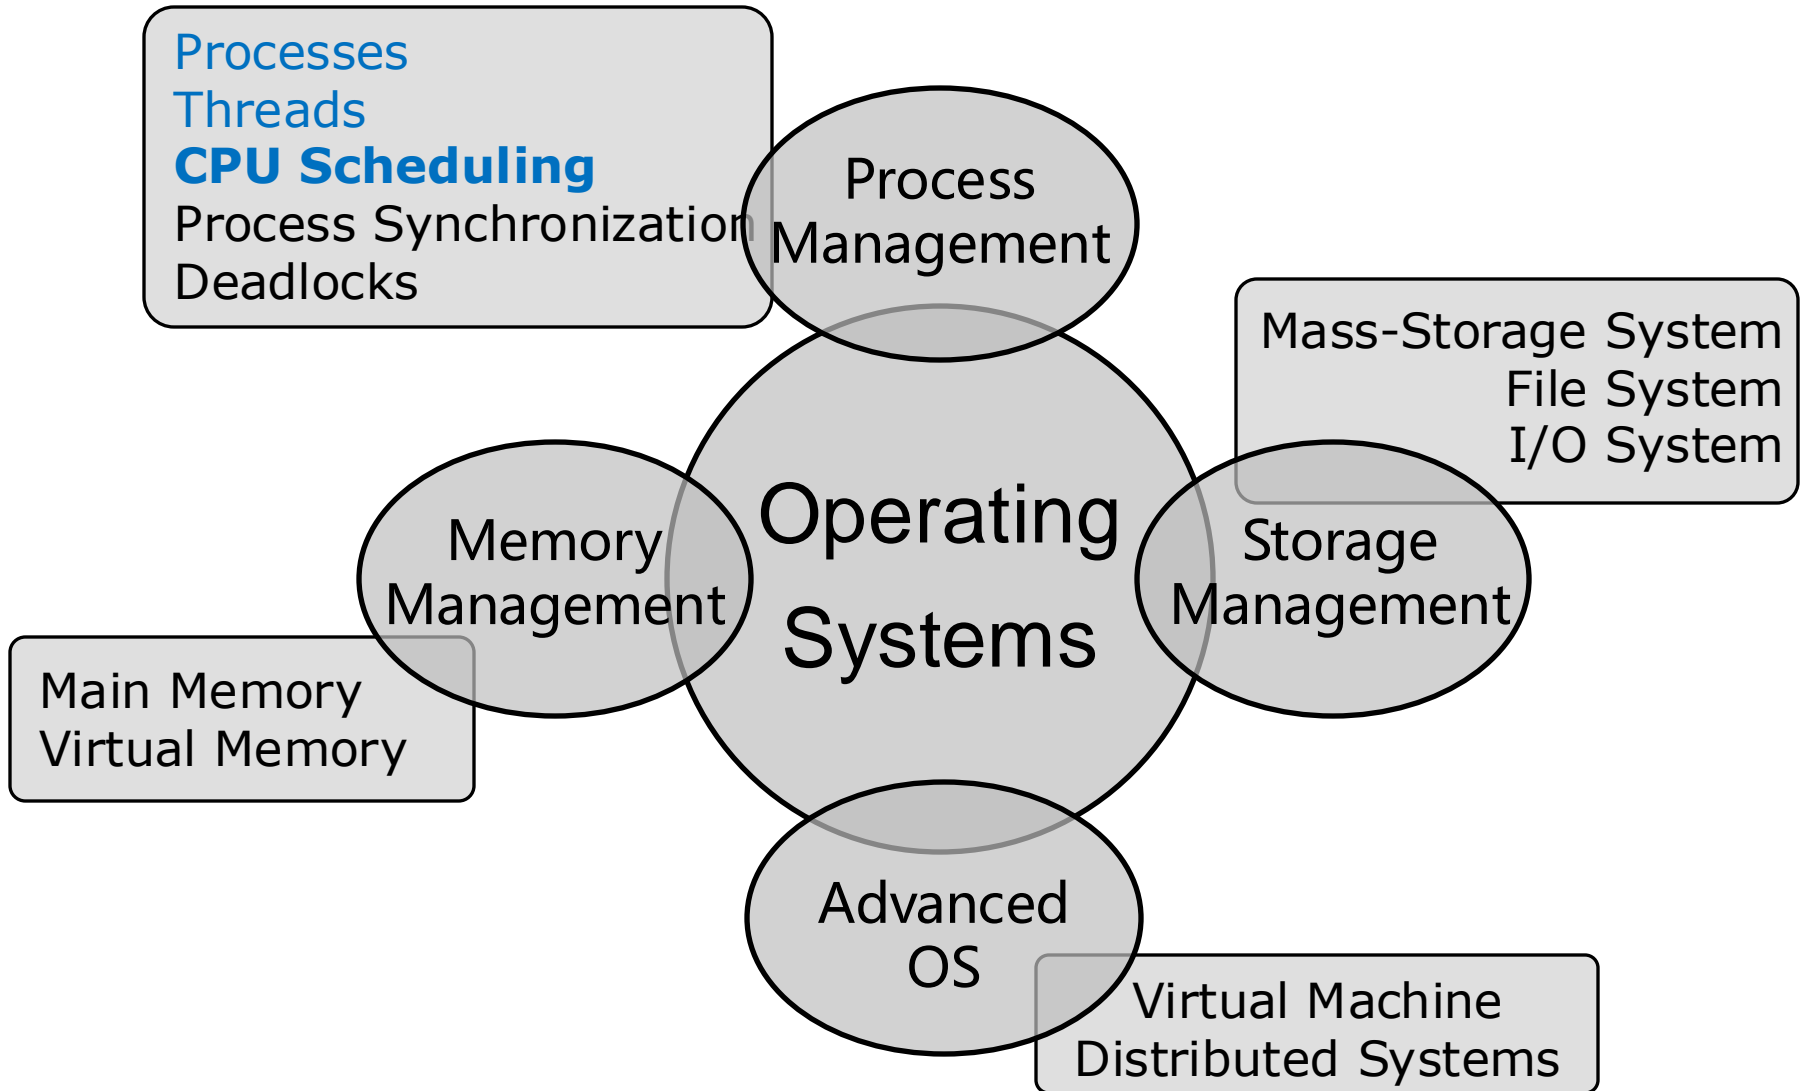


CPU Scheduling

Shengzhong Liu

Department of Computer Science and Engineering
Shanghai Jiao Tong University

Operating System Topics



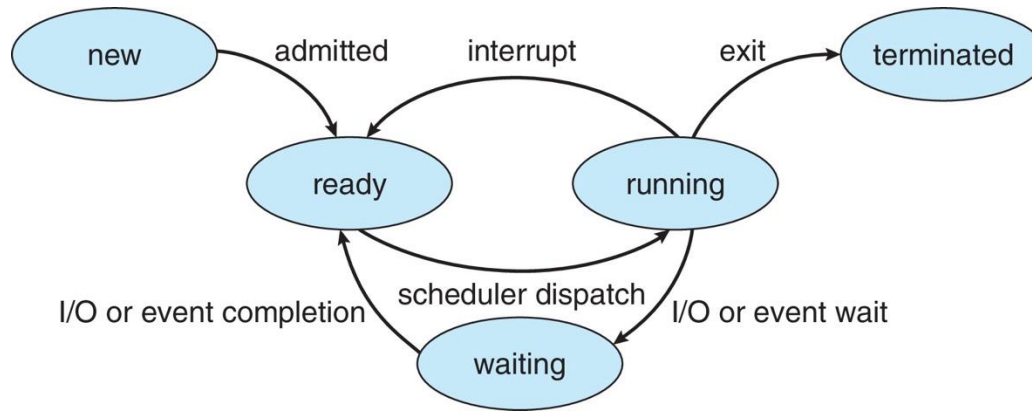
Outline

- Basic Concepts
- Scheduling Algorithms
- Thread Scheduling
- Multi-Processor Scheduling
- Real-Time CPU Scheduling
- Algorithm Evaluation: Deterministic Modeling

Basic Concept


Process State and Queue

Process state transition




Ready queue and wait queue

- Ready queue – processes wait to execute
- Wait queue – processes waiting for an event (i.e., I/O)




Ready Queue



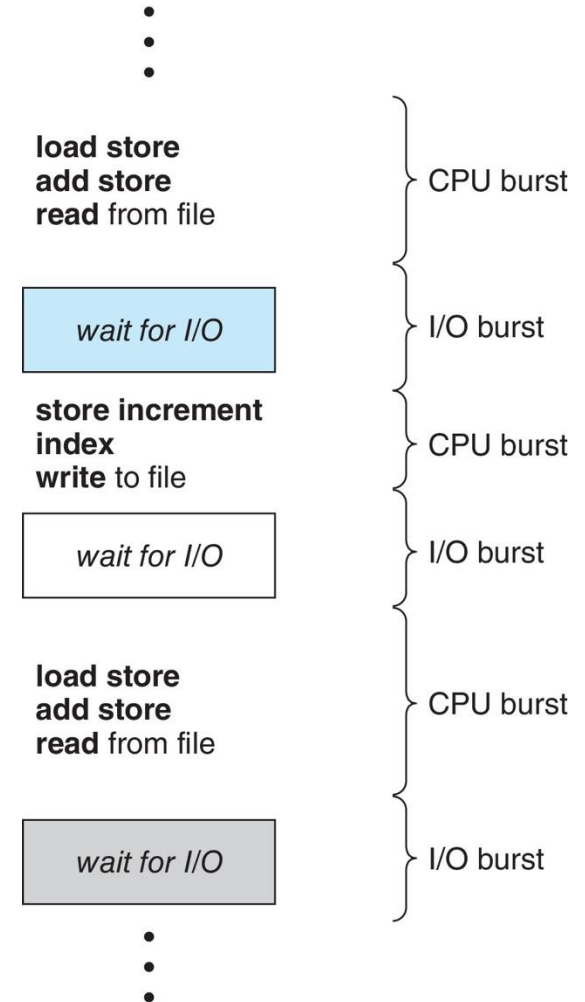
Wait Queue

Basic Concepts

- ❑ Maximize CPU utilization obtained with multiprogramming
- ❑ CPU-I/O Burst Cycle
 - ❑ Process execution consists of an alternating sequence of **CPU bursts** and **I/O bursts**
- ❑ CPU burst distribution is of main concern



Histogram of CPU-burst durations



CPU Scheduler

- Selects from the processes in ready queue, and allocates the CPU to one
 - Queue may be ordered in various ways
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready state
 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**, scheduling under 2 and 3 is **preemptive**
 - **Nonpreemptive Scheduling**: Once the CPU has been allocated to a process, the process keeps the CPU until it releases the CPU either by terminating or by switching to the waiting state
 - **Preemptive Scheduling**: The CPU is allocated to the process for a limited amount of time and then taken away

Preemptive and Nonpreemptive Scheduling

□ Preemptive scheduling

vs.

Non-preemptive scheduling

Process	Arrival Time	CPU Burst Time (in millisec.)
P0	3	2
P1	2	4
P2	0	6
P3	1	4



Preemptive Scheduling

Process	Arrival Time	CPU Burst Time (in millisec.)
P0	3	2
P1	2	4
P2	0	6
P3	1	4




Non-Preemptive Scheduling

	Nonpreemptive	Preemptive
CPU Utilization	Lower	Higher
Response Time	Longer	Shorter
Complexity	Lower	Higher
Priority Management	Static	Allows dynamic priority
Fairness	Worse	Better

Dispatcher (分发器)

- ❑ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler. This involves:
 - ❑ **Switching context**
 - ❑ Switching to user mode
 - ❑ Jumping to the proper location in the user program to restart that program
- ❑ **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running



Scheduling Criteria / Metrics

- **CPU utilization** – keep the CPU as busy as possible
 - Max CPU utilization
- **Throughput** – # of processes that complete execution per time unit
 - Max throughput
- **Turnaround time** – amount of time to execute a particular process.
 - The interval from the time of submission of a process to the time of completion is the turnaround time.
 - Min turnaround time
- **Waiting time** – amount of time a process has been waiting in the ready queue
 - Min waiting time
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced
 - Min response time

Scheduling Algorithms


CPU Scheduling Algorithms

- ❑ First-Come, First-Served Scheduling (**FCFS**)
- ❑ Shortest-Job-First Scheduling (**SJF**)
- ❑ Priority Scheduling (**PS**)
- ❑ Round-Robin Scheduling (**RR**)
- ❑ Multilevel Queue Scheduling (**MQS**)
- ❑ Multilevel Feedback Queue Scheduling (**MFQS**)

1. First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>CPU Burst</u>	<u>Arrival Time</u>
P_1	24	0
P_2	3	1
P_3	3	2

□ The **Gantt Chart** for the schedule is:



□ **Waiting time** for $P_1 = 0$; $P_2 = 24 - 1 = 23$; $P_3 = 27 - 2 = 25$

Average waiting time: $(0 + 23 + 25)/3 = 16$


□ **Turnaround time** for $P_1 = 24$; $P_2 = 27 - 1 = 26$; $P_3 = 30 - 2 = 28$

Average turnaround time: $(24 + 26 + 28)/3 = 26$

FCFS Scheduling (Cont.)

<u>Process</u>	<u>CPU Burst</u>	<u>Arrival Time</u>
P_1	24	2
P_2	3	0
P_3	3	1

- The **Gantt chart** for the schedule is:




- Waiting time for $P_1 = 4$; $P_2 = 0$; $P_3 = 2$
- Average waiting time: $(4 + 0 + 2)/3 = 2$
- Much better than previous case
- **Convoy effect** - short process behind long process
 - Consider one CPU-bound and many I/O-bound processes

FCFS Scheduling (Cont.)

Process	CPU Burst	I/O Burst	CPU Burst	Arrival Time
P_1	12	3	12	0
P_2	1	2	2	1
P_3	1	2	2	2

□ The Gantt Chart for the schedule is:



□ Waiting time for

□ $P_1 = 15 - 12 - 3 = 0$

□ $P_2 = (12 - 1) + (27 - 13 - 2) = 23$

□ $P_3 = (13 - 2) + (29 - 14 - 2) = 24$

□ Turnaround time for $P_1 = 27$; $P_2 = 29 - 1 = 28$; $P_3 = 31 - 2 = 29$

□ CPU utilization $30/31 = 96.77\%$


2. Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of **its next CPU burst**
 - Use these lengths to schedule the process with the shortest time
- SJF is **optimal** – gives minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request
- Preemptive version called **shortest-remaining-time-first**

Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

- SJF scheduling chart (assume all processes arrive at 0)



- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

Determining Length of Next CPU Burst

- Can only estimate the length – should be similar to the previous one
 - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, using **exponential moving average**

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n.$$

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$

- Commonly, α is set to $\frac{1}{2}$

Examples of Exponential Averaging

□ $\alpha = 0$

□ $\tau_{n+1} = \tau_n$

□ Recent history does not count

□ $\alpha = 1$

□ $\tau_{n+1} = t_n$


□ Only the actual last CPU burst counts

□ If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$

□ Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

Prediction of the Length of the Next CPU Burst



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

Shortest Remaining Time First Scheduling


- Preemptive version of SJF
- Whenever a new process arrives in the ready queue, the decision on which process to schedule next is redone using the SJF algorithm.
- Is SRT more “optimal” than SJF in terms of the **minimum average waiting time** for a given set of processes?

Example of Shortest-remaining-time-first

- We now add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- **Preemptive** SJF Gantt Chart



- Average waiting time = $[(10-1)+(1-1)+(17-2)+(5-3)]/4 = 26/4 = 6.5$


3. Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - Nonpreemptive
- SJF is priority scheduling, where priority is the inverse of predicted next CPU burst time
- Problem: **Starvation** – low priority processes may never execute
- Solution: **Aging** – as time progresses increase the priority of the process

Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

□ Priority scheduling Gantt Chart



□ Average waiting time = 8.2 msec


4. Round Robin Scheduling (RR)

- Round Robin (RR) is similar to FCFS scheduling, but preemption is added to switch between processes.
- Each process gets a small unit of CPU time (**time quantum** q), usually 10-100 milliseconds.
 - After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time of at most q time units at once.
 - No process waits more than $(n-1)q$ time units before next execution.
- Timer interrupts every quantum to schedule next process
- Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3


□ The Gantt chart is:




- Typically, higher average turnaround than SJF, but better *response*
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch $< 10 \mu s$
- What's the number of context switch?

Time Quantum and Context Switch Time

□ Context switching



Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7


Rule of thumb:
80% of CPU bursts should
be shorter than quantum

5. Multilevel Queue Scheduling

- Ready queue is partitioned into separate queues, e.g.:
 - foreground (interactive)
 - background (batch)
- Process joins a given queue
- Each queue has its own scheduling algorithm:
 - foreground – RR
 - background – FCFS
- Scheduling must be done between the queues:
 - Fixed priority scheduling: (i.e., serve all from foreground then from background).
 - ▶ Possibility of **starvation**.
 - Time slice – each queue gets a certain amount of CPU time, which it can schedule amongst its processes, e.g., 80% to foreground in RR, 20% to background in FCFS

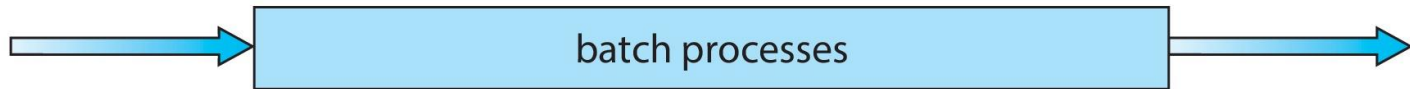
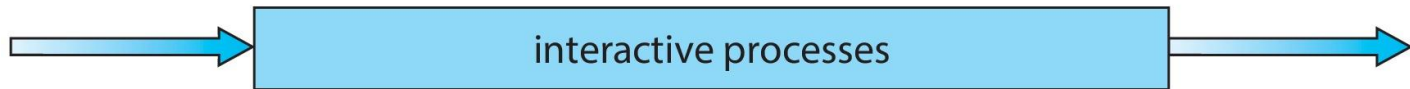

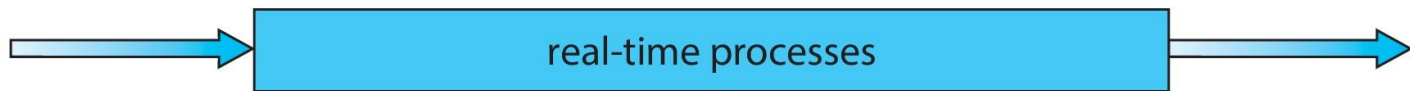
Multilevel Queue

- With priority scheduling, have separate queues for each priority.
- Schedule the process in the highest-priority queue!



Multilevel Queue Scheduling

highest priority




lowest priority

6. Multilevel Feedback Queue Scheduling

- A process can move between the various queues
 - Aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - Queue change:
 - ▶ method used to determine when to upgrade a process
 - ▶ method used to determine when to downgrade a process
 - ▶ method used to determine which queue a process will enter when that process needs service
- Aging can be implemented using multilevel feedback queue

Example of Multilevel Feedback Queue

- Three queues:
 - Q_1 – RR with time quantum 8 milliseconds
 - Q_2 – RR with time quantum 16 milliseconds
 - Q_3 – FCFS



Multilevel Feedback Queue


□ Scheduling

- A new job enters queue Q_1 which is served FCFS
 - ▶ When it gains CPU, job receives 8 milliseconds
 - ▶ If it does not finish in 8 milliseconds, job is moved to queue Q_2
- At Q_2 job is again served FCFS/RR and receives 16 additional milliseconds
 - ▶ If it still does not complete, it is preempted and moved to queue Q_3
- If a process does not use up its quantum in the current level, it will keep its current queuing level and be put into the end of the queue.
 - ▶ Then, it can still get the same amount of quantum (not remaining quantum) next time when it is picked.

Example of Using Multilevel Feedback Queue

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	36
P_2	16	20
P_3	20	12


□ The Gantt chart is:



Example of Using Multilevel Feedback Queue

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	36
P_2	15	20
P_3	20	12


□ The Gantt chart is:



Thread Scheduling

Thread Scheduling

- Distinction between user-level and kernel-level threads
- When threads are supported, threads are scheduled instead of processes
- In many-to-one and many-to-many models, the thread library schedules user-level threads to run on kernel-level threads
 - Known as **process-contention scope (PCS, 进程竞争范围)** since scheduling competition happens within the process
 - Typically done via priority set by the programmer
- Kernel thread scheduled onto available CPU is **system-contention scope (SCS, 系统竞争范围)** – competition among all threads in the system
 - One-to-one mapping use only SCS.



Pthread Scheduling

- API allows specifying either PCS or SCS during thread creation
 - **PTHREAD_SCOPE_PROCESS** schedules threads using PCS scheduling
 - ▶ Schedules user-level threads onto available LWPs
 - ▶ Number of LWPs is maintained by the thread library
 - **PTHREAD_SCOPE_SYSTEM** schedules threads using SCS scheduling
 - ▶ Creates and binds an LWP for each user-level thread
 - ▶ In fact, implements the one-to-one mapping
- Can be limited by OS – Linux and Mac OS X only allow PTHREAD_SCOPE_SYSTEM

Pthread Scheduling API

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
int main(int argc, char *argv[]) {
    int i, scope;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* first inquire on the current scope */
    if (pthread_attr_getscope(&attr, &scope) != 0)
        fprintf(stderr, "Unable to get scheduling scope\n");
    else {
        if (scope == PTHREAD_SCOPE_PROCESS)
            printf("PTHREAD_SCOPE_PROCESS");
        else if (scope == PTHREAD_SCOPE_SYSTEM)
            printf("PTHREAD_SCOPE_SYSTEM");
        else
            fprintf(stderr, "Illegal scope value.\n");
    }
}
```


Pthread Scheduling API

```
/* set the scheduling algorithm to PCS or SCS */
pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
/* create the threads */
for (i = 0; i < NUM_THREADS; i++)
    pthread_create(&tid[i], &attr, runner, NULL);
/* now join on each thread */
for (i = 0; i < NUM_THREADS; i++)
    pthread_join(tid[i], NULL);
}

/* Each thread will begin control in this function */
void *runner(void *param)
{
    /* do some work ... */
    pthread_exit(0);
}
```


Multi-Processor Scheduling

Multiple-Processor Scheduling

- CPU scheduling is more complex when multiple CPUs are available
 - **Homogeneous processors** within a multiprocessor

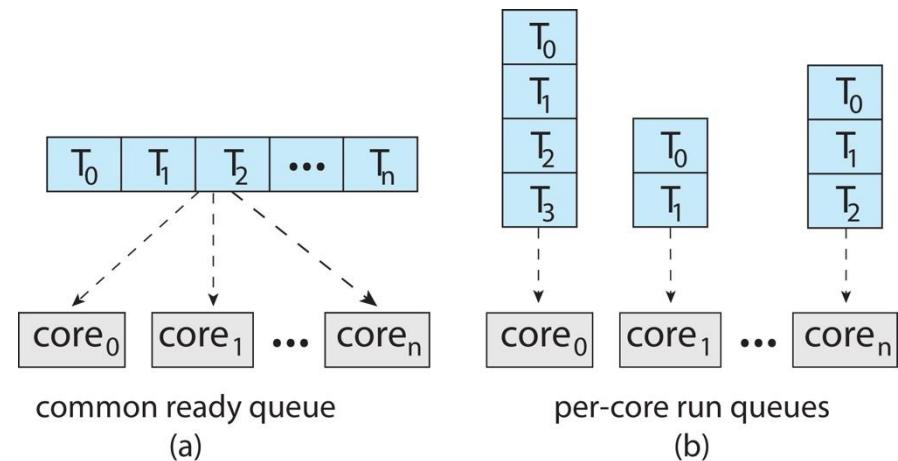
- **Asymmetric multiprocessing**

- One processor handles all scheduling decisions, I/O processing, and other system activities
- Only one processor accesses the system data structures, alleviating the need for data sharing
- All other processors just execute user code.



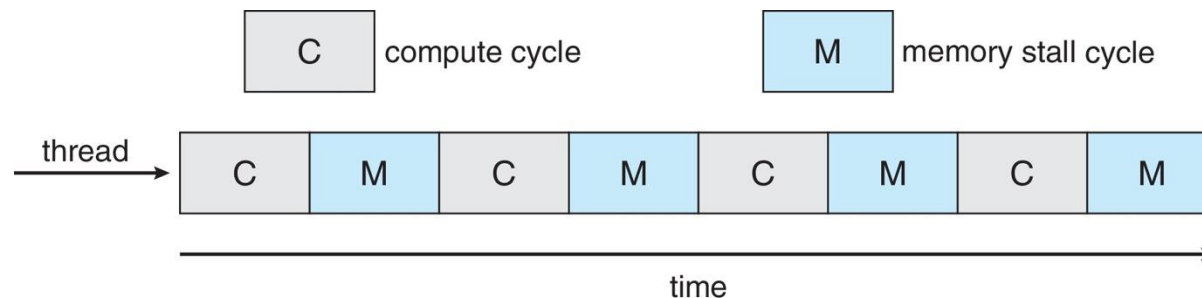
- **Symmetric multiprocessing (SMP)**

- Each processor is self-scheduling,
- (a) All processes are in common ready queue,
- (b) Each processor has its own private queue of ready processes




Multicore Processors

- Recent trend to place multiple processor cores on same physical chip
 - Faster and consumes less power
- **Memory stall:**
 - When a processor accesses memory, it spends a significant amount of time waiting for the data to become available
- Multiple threads per core also growing
 - Takes advantage of memory stall to make progress on another thread while memory retrieve happens



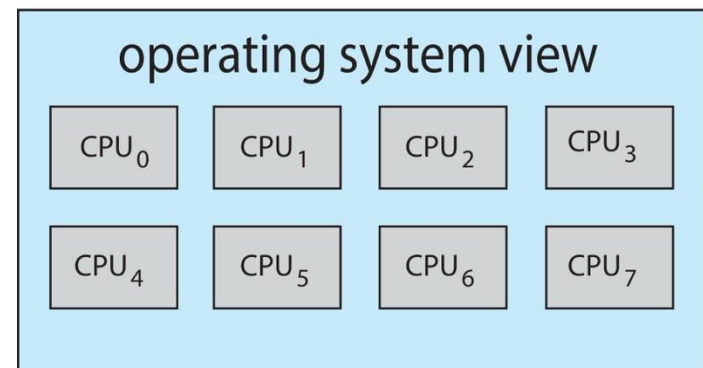
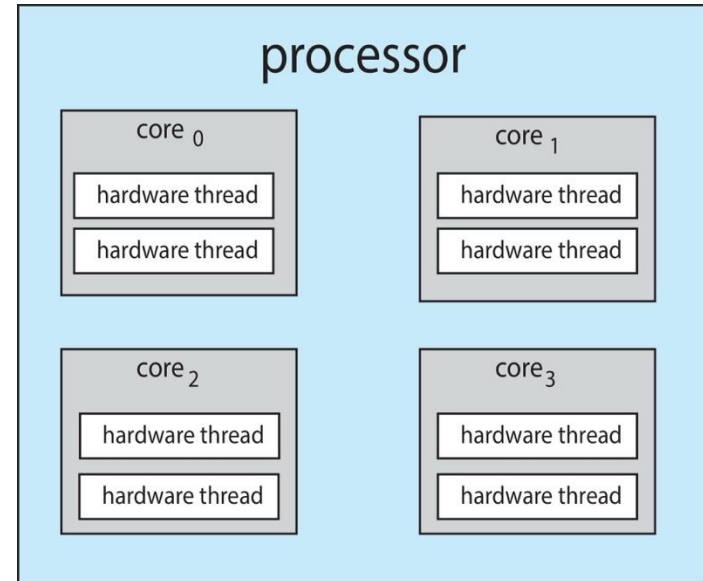
Multithreaded Multicore System

- ❑ Each core has > 1 hardware threads.
- ❑ If one thread has a memory stall, switch to another thread!
- ❑ Figure



Multithreaded Multicore System

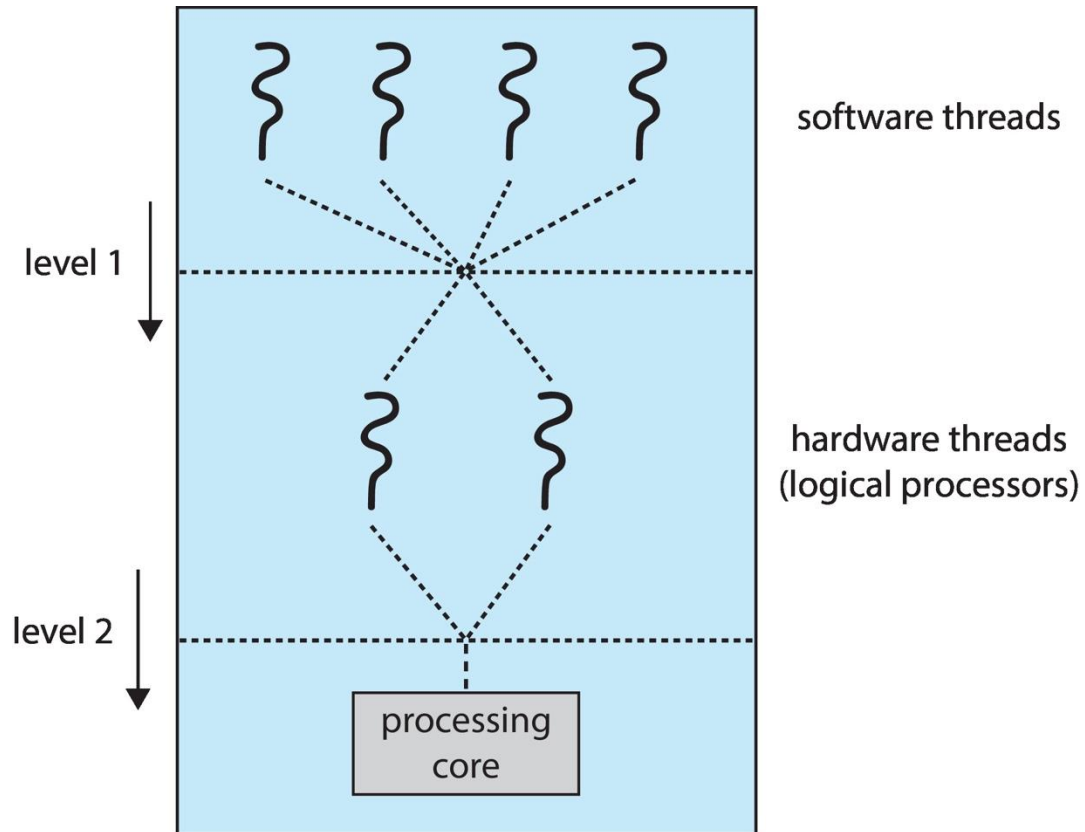
- ❑ **Chip-multithreading (CMT)** assigns each core multiple hardware threads.
 - ❑ Intel refers to this as **hyperthreading**.
- ❑ On a quad-core system with 2 hardware threads per core, the operating system sees 8 logical processors.



Multithreaded Multicore System

□ Two levels of scheduling:

1. The operating system decides which **software thread** to run on which **hardware thread**
2. Each core decides which **hardware thread** to run on the **physical core**



Multiple-Processor Scheduling – Load Balancing

- If SMP, need to keep all CPUs loaded for efficiency
- **Load balancing** attempts to keep workload evenly distributed
- Two types of migration:
 - **Push migration** – periodic task checks load on each processor, and if found overloading, pushes task from overloaded CPU to other CPUs
 - **Pull migration** – idle processors pulls waiting task from busy processor


Multiple-Processor Scheduling – Processor Affinity

□ Processor affinity:

- When a thread has been running on one processor, the cache contents of that processor stores the memory accesses by that thread.
- We refer to this as a thread having affinity for a processor
- Load balancing may affect processor affinity as a thread may be moved from one processor to another to balance loads
 - The thread loses the contents in the cache of the previous processor.

□ Two types:

- **Soft affinity** – the operating system attempts to keep a thread running on the same processor, but no guarantees.
- **Hard affinity** – allows a process to specify a set of processors it may run on.




Real-Time CPU Scheduling

Real-Time CPU Scheduling

- Real-time scheduling allocates CPU resources to tasks based on their timing constraints, ensuring they meet strict deadlines for predictable and deterministic execution.
- Can present obvious challenges
- Two types:
 - **Soft real-time systems** – Critical real-time tasks have the highest priority, but no guarantee as to when tasks will be scheduled
 - **Hard real-time systems** – Task must be serviced by its deadline


Priority-based Scheduling

- For real-time scheduling, scheduler must support **preemptive, priority-based** scheduling
 - But only guarantees soft real-time
- For hard real-time must also provide ability to meet deadlines
- Processes have new characteristics: **periodic** ones require CPU at constant intervals
 - Has processing time t , deadline d , period p
 - $0 \leq t \leq d \leq p$
 - **Rate** of periodic task is $1/p$




RT Scheduling: Rate Monotonic

- A priority is assigned based on the inverse of its period
- Shorter periods = higher priority;
- Longer periods = lower priority
- P_1 is assigned a higher priority than P_2 .




- Process P_2 misses finishing its deadline at time 80



RT Scheduling: Earliest Deadline First (EDF)

- Priorities are assigned according to deadlines:
 - The earlier the deadline, the higher the priority
 - The later the deadline, the lower the priority

□ Figure



Algorithm Evaluation


Algorithm Evaluation

- How to select a CPU-scheduling algorithm for an OS?
 - Determine criteria, then evaluate algorithms
- Analytical evaluation:
 - Use the given algorithm and the system workload to produce a formula or number to evaluate the performance of the algorithm for the workload.
- **Deterministic modeling**, one type of analytic evaluation:
 - Takes a particular predetermined workload and defines the performance of each algorithm for that workload
- Consider 5 processes arriving at time 0:

<u>Process</u>	<u>Burst Time</u>
P_1	10
P_2	29
P_3	3
P_4	7
P_5	12


Deterministic Evaluation

- For each algorithm, calculate minimum **average waiting time**
- Simple and fast, but requires exact numbers for input, applies only to those inputs
 - FCFS is **28ms**:




Process	Burst Time
P_1	10
P_2	29
P_3	3
P_4	7
P_5	12

- Non-preemptive SFJ is **13ms**:



- RR is **23ms**:



Summary

- ❑ CPU scheduling is the task of selecting a waiting process from the ready queue and allocating the CPU to it.
- ❑ Scheduling algorithms may be either preemptive or nonpreemptive. Most modern operating systems are preemptive.
- ❑ Scheduling algorithm criteria: (1) CPU utilization, (2) throughput, (3) turnaround time, (4) waiting time, and (5) response time.
- ❑ Introduced scheduling algorithms: FCFS, SJF, RR, priority scheduling, multilevel queue scheduling, multilevel feedback queue scheduling
- ❑ Multicore processors place one or more CPUs on the same physical chip, and each CPU may have more than one hardware thread.
- ❑ Real-time scheduling provides timing guarantees for real-time tasks

Homework

- Reading
 - Chapter 5
- Homework
 - Check Canvas for Homework 1 (later today), due on March 12 at 23:59.
 - Write your homework with the given LaTeX template.