

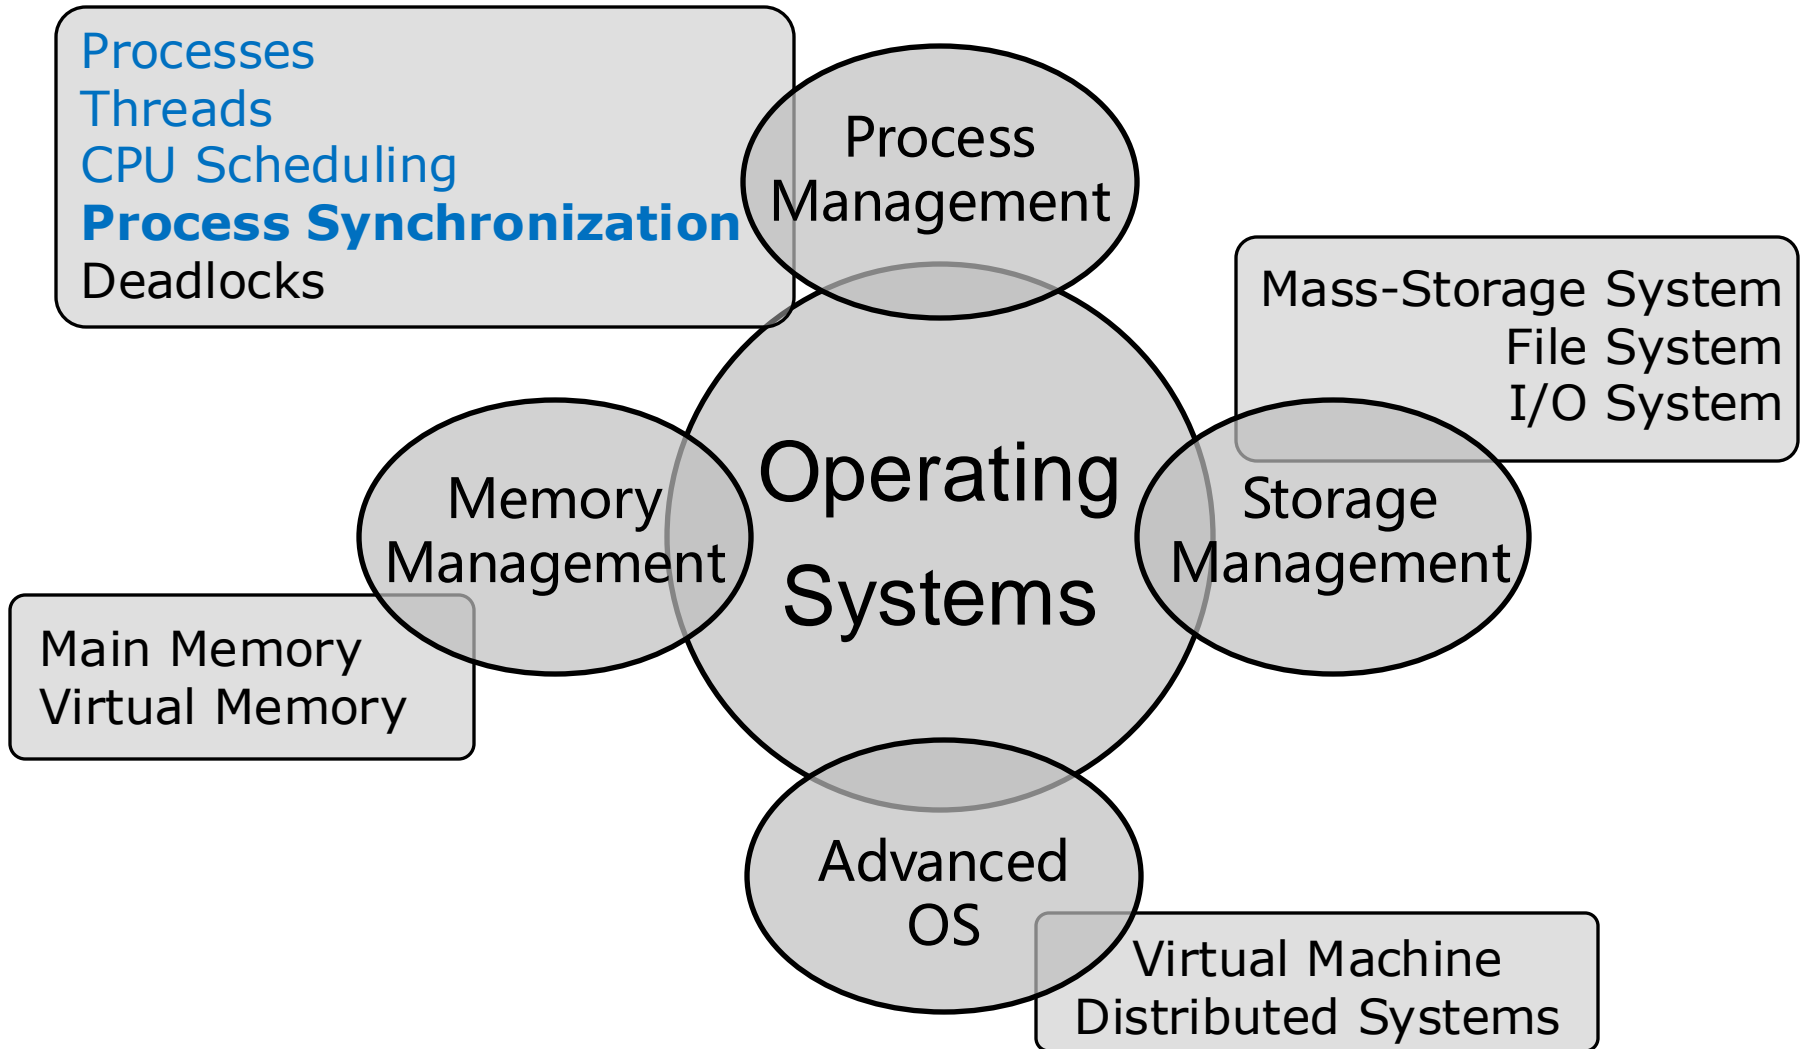
# Process Synchronization

Shengzhong Liu

Department of Computer Science and Engineering

Shanghai Jiao Tong University

# Operating System Topics



# Outline


- Background: Critical-Section Problem
- Synchronization Mechanisms:
  - Peterson's Solution
  - Hardware Support for Synchronization
  - Mutex Locks and Semaphores
  - Monitors
- Synchronization Problem Formulations

# Background:

# Critical Section Problem


# Background

- Processes can execute concurrently
  - May be interrupted at any time, partially completing execution
- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the **orderly execution of cooperating processes**



# Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
  - **Unbounded-buffer** places no practical limit on the buffer size
  - **Bounded-buffer** assumes that there is a fixed buffer size



# Improved Solution of Producer-Consumer Problem

<pre>while (true) {     /* produce an item and put in nextProduced*/     while (<b>counter == BUFFER_SIZE</b>); // do nothing     buffer [in] = nextProduced;     in = (in + 1) % BUFFER_SIZE;     <b>counter++</b>; }</pre>	<b>Producer</b>
--	-----------------

This solution allows us to utilize all available buffer slots.

<pre>while (true) {     while (<b>counter == 0</b>) ; // do nothing     nextConsumed = buffer[out];     out = (out + 1) % BUFFER_SIZE;     <b>counter--</b>;     /* consume the item */ }</pre>	<b>Consumer</b>
---	-----------------

# Race Condition: Example 1

- `counter++` could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- `counter--` could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution interleaving with “`counter = 5`” initially:

S0: producer execute	<code>register1 = counter</code>	{register1 = 5}
S1: producer execute	<code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute	<code>register2 = counter</code>	{register2 = 5}
S3: consumer execute	<code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute	<code>counter = register1</code>	{counter = 6}
S5: consumer execute	<code>counter = register2</code>	{counter = 4}

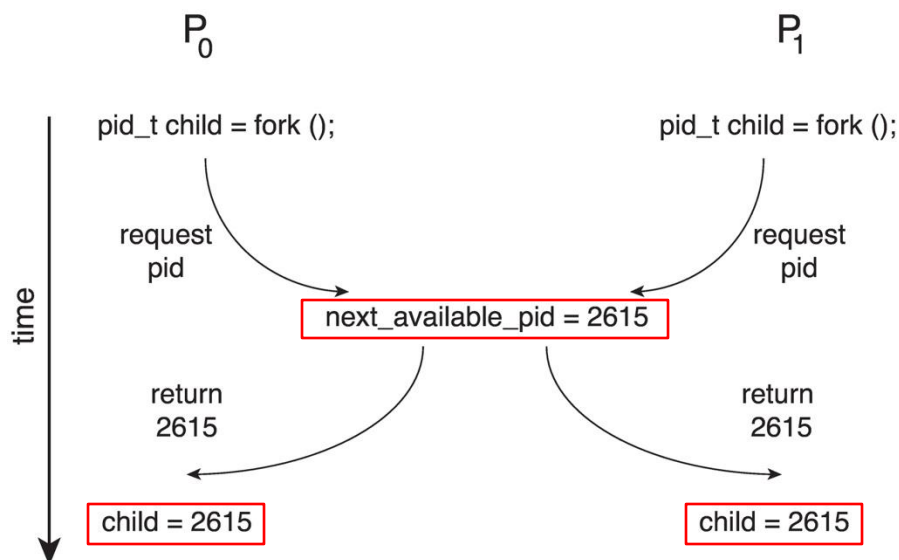
- **Race condition:**

- Several processes access and manipulate the same data concurrently
- The execution outcome depends on access order



# Race Condition: Example 2

- Processes  $P_0$  and  $P_1$  are creating child processes using the `fork()` system call
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (`pid`)



- Unless there is a mechanism to prevent  $P_0$  and  $P_1$  from accessing the variable `next_available_pid`, the same pid could be assigned to two different processes!

# Critical Section Problem

- Consider system of  $n$  processes  $\{p_0, p_1, \dots, p_{n-1}\}$ 
    - Each process has a **critical section** segment of codes
    - Process may be changing common variables, updating table, writing file, etc
    - When one process is in critical section, no others could in its critical section
  - Critical section problem is to design protocols to solve this
  - Each process must:
    - ask permission to enter the critical section in **entry section**
    - may follow critical section with **exit section**
    - then **remainder section**
  - Especially challenging with preemptive kernels
- ```
while (true) {  
    entry section  
    critical section  
    exit section  
    remainder section  
}
```

# Solution Requirements of CS Problem

## 1. Mutual Exclusion


- If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections

## 2. Progress

- If no process is executing in its critical section and there exist processes waiting to enter, the selection of entering processes cannot be postponed indefinitely

## 3. Bounded Waiting


- A bound must exist on #times that other processes enter their critical sections after a process requests to enter its critical section but before the request is granted
- **Sequential Access** – The sequence of accessing the critical section follows the order of requests raised by the processes



# Synchronization Mechanisms

# Process Synchronization Mechanisms

- ❑ Peterson's Solution – A classical and general algorithm
- ❑ Hardware Synchronization
- ❑ Mutex and Semaphores
- ❑ Monitors



# Peterson's Solution

- We have two processes:  $P_i$  and  $P_j$
- Assume the **load** and **store** machine-language instructions are atomic;
  - They cannot be interrupted
- The two processes share two variables:
  - **int** `turn`;
  - **boolean** `flag[2]`
- The variable **turn** indicates whose turn it is to enter the critical section
- The **flag** array is used to indicate if a process is ready to enter the critical section.
  - `flag[i] = true` implies that process  $P_i$  is ready!

```
while (true) {
```

```
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j)
        ;
```

```
    /* critical section */
```

```
    flag[i] = false;
```

```
    /* remainder section */
```

```
}
```

Algorithm for **process  $P_i$**

# Correctness of Peterson's Solution

□ Provable that the three CS requirement are met:

1. Mutual exclusion is preserved

$P_i$  enters CS only if: either `flag[j] = false` or `turn = i`

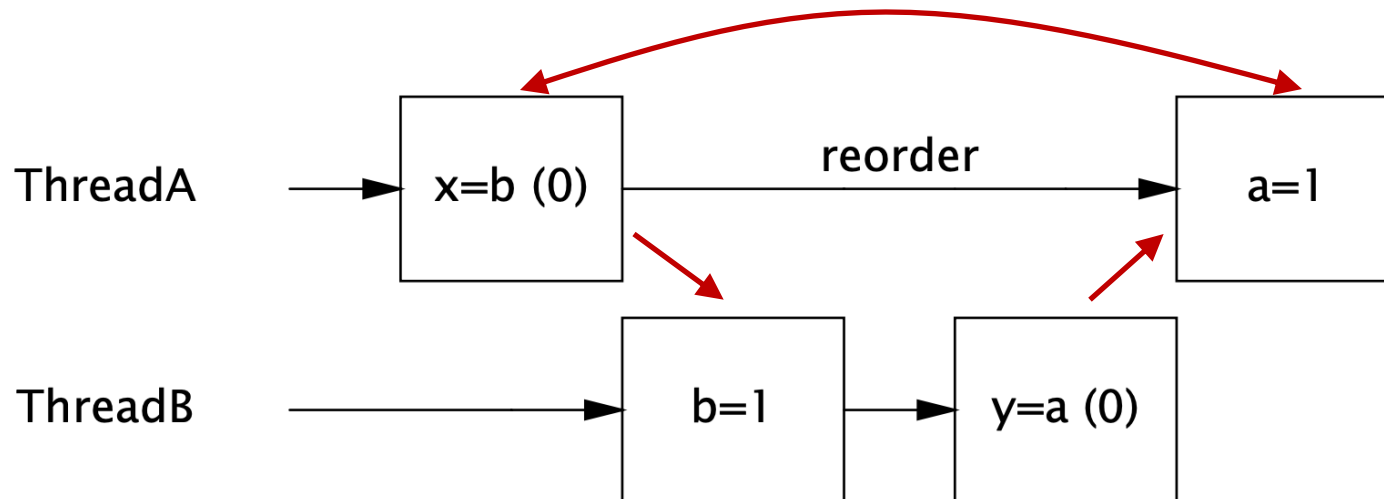
2. Progress requirement is satisfied

3. Bounded-waiting requirement is met

```
while (true){  
  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
  
    /* critical section */  
  
    flag[i] = false;  
  
    /* remainder section */  
  
}
```

# Peterson's Solution and Modern Architecture

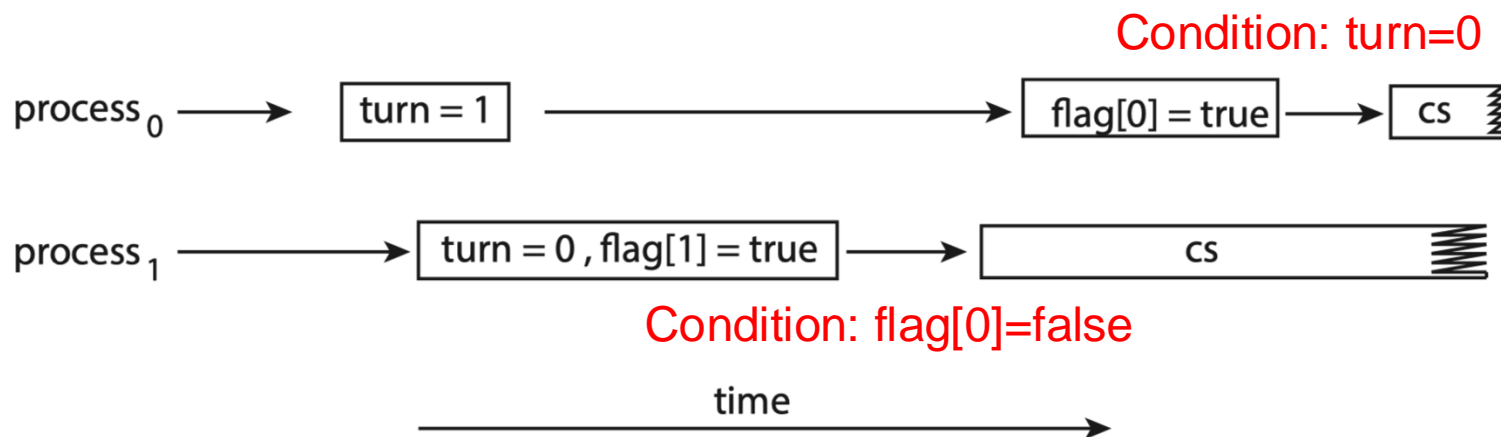
- Although useful for demonstrating an algorithm, Peterson's Solution is not guaranteed to work on modern architectures.
  - To improve performance, processors and/or compilers may **reorder read and write operations that have no dependencies**
  - For single-threaded program, this is ok as the result will always be the same.
  - For multi-threaded program, the reordering may produce inconsistent or unexpected results!





# Peterson's Solution

- Peterson's Solution is not guaranteed to work on modern architectures with reordered execution of instructions.
- We need proper synchronization tools to preserve mutual exclusion.
- The example below allows both processes to be in their critical section at the same time!



# Synchronization Mechanisms: Hardware Synchronization

# Hardware Support for Synchronization

- Many systems provide hardware support for critical section code
  - Memory barriers
  - Hardware instructions
  - Atomic variables

# Hardware Sync: (1) Memory Barrier

- **Memory models** are the memory guarantees a computer architecture makes to application programs, including two categories:
  - **Strongly ordered** – a memory modification of one processor is **immediately visible** to all other processors.
  - **Weakly ordered** – a memory modification of one processor **may not be immediately visible** to all other processors.
- A **memory barrier (内存屏障)** is an instruction that forces any change in memory to be propagated (visible) to all other processors.
  - When a memory barrier instruction is performed, **all previous loads and stores are completed before any subsequent load / store operations are performed (within the same process)**.
- Even if instructions were reordered, the memory barrier ensures that the store operations are
  - Completed in memory
  - Visible to other processors before future load / store operations

# Reordered Execution Example

- Two threads share the data:

```
boolean flag = false;  
int x = 0;
```

- Thread 1 performs

```
while (!flag)  
;  
print x
```

- Thread 2 performs

```
x = 100;  
flag = true
```

- What is the expected output?

**100**

- However, since the variables `flag` and `x` are independent of each other, the instructions:

```
flag = true;  
x = 100;
```

for Thread 2 may be reordered

- If this occurs, the output may be **0!**

# Hardware Sync: (1) Memory Barrier

- Memory barrier:
  - All loads and stores are completed before any subsequent load or store operations are performed.
  - We could add a memory barrier to ensure Thread 1 outputs 100
- Thread 1 now performs

```
while (!flag)
    memory_barrier();
print x
```
- Thread 2 now performs

```
x = 100;
memory_barrier();
flag = true
```
- For Thread 1 we guarantee that
  - The value of `flag` is loaded before the value of `x`.
- For Thread 2 we guarantee that
  - The assignment to `x` occurs before the assignment `flag`.

# Hardware Sync: (2) Hardware Instructions

- Atomic instruction:
  - The instruction can not be interrupted by other instructions.
  - Appears to occur instantaneously from the view of other threads or processes executing concurrently.
  - Either completes entirely or has no effect at all
- Two abstract atomic instruction types:
  - **Test-and-Set** instruction
  - **Compare-and-Swap** instruction

# The `test_and_set` Instruction

- Definition of `test_and_set` instruction:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = true;
    return rv;
}
```

- Properties

- Executed atomically
- Set the new value of passed parameter to `true`
- Returns the original value of passed parameter



# The `compare_and_swap` Instruction

- Definition of `compare_and_swap` instruction:

```
int compare_and_swap(int *value, int expected, int new_value)
{
    int temp = *value;
    if (*value == expected)
        *value = new_value;
    return temp;
}
```

- Properties

- Executed atomically
- Returns the original value of the passed parameter `value`
- Set the variable `value` as the passed parameter `new_value` but only if `*value == expected` is true.
  - The swap happens only under this condition.

# Hardware Instructions for Mutual Exclusion

## test\_and\_set solution

- Shared boolean variable `lock`, initialized to `false`
- Solution:

```
do {
    while (test_and_set(&lock))
        ; /* do nothing */

    /* critical section */

    lock = false;
    /* remainder section */
} while (true);
```

## compare\_and\_swap solution

- Shared integer `lock` initialized to 0;
- Solution:

```
while(true){
    while (compare_and_swap(&lock,
0, 1) != 0)
        ; /* do nothing */

    /* critical section */

    lock = 0;
    /* remainder section */
}
```

Do they solve the critical-section problem?

# Bounded-waiting with `compare_and_swap`

```
while (true) {
    /* entry section - acquire the lock */
    waiting[i] = true;
    key = 1;
    while (waiting[i] && key == 1) // use key to exit
        key = compare_and_swap(&lock, 0, 1);
    waiting[i] = false;
    /* critical section - release one process */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = 0;
    else
        waiting[j] = false;
    /* remainder section */
}
```

# Hardware Sync: (3) Atomic Variables

- Typically, instructions such as compare-and-swap are used as building blocks for other synchronization tools.
- One way to use is defining **atomic variable** that provides atomic updates (uninterruptible) on basic data types such as integers and booleans.
- For example:
  - Let **sequence** be an atomic variable
  - Let **increment()** be operation on the atomic variable **sequence**
  - The command:

**increment(&sequence) ;**

ensures **sequence** is incremented without interruption:

# Hardware Sync: (3) Atomic Variables

- The `increment()` function can be implemented with `compare_and_swap` instruction as follows:

```
void increment(atomic_int *v)
{
    int temp;
    do {
        temp = *v;
    }
    while (temp != (compare_and_swap(v, temp, temp+1)));
}
```

Goal:  $v \rightarrow v+1$

# Synchronization Mechanisms: Mutex and Semaphore

# Mutex Locks (互斥锁)

- OS designers build higher-level software tools to solve critical section problem, where the simplest is **mutex** lock

- Mutex is a boolean variable indicating if lock is available or not

- Protect a critical section by

- First **acquire()** a lock
  - Then **release()** the lock


- Calls to **acquire()** and **release()** must be **atomic**

- Usually implemented via hardware atomic instructions such as compare-and-swap.

- But this solution requires **busy waiting**

- This lock therefore called a **spinlock (自旋锁)**

```
while (true) {  
    acquire lock  
    critical section  
    release lock  
remainder section  
}
```



# Semaphore (信号量)

- Semaphore is a synchronization tool that provides more sophisticated ways (than Mutex locks) for processes to synchronize their activities.
  - Semaphore **S** – integer variable
- Can only be accessed via two **atomic operations**
  - **wait()** and **signal()**
    - ▶ Originally called **P()** and **V()**
- Definition of the **wait()** operation      □ Definition of the **signal()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```



# Semaphore (Cont.)

- Two types of semaphore:
  - **Counting semaphore**
    - ▶ Integer value can range over an unrestricted domain
  - **Binary semaphore**
    - ▶ Integer value can range only between 0 and 1
    - ▶ Same as a **mutex lock**
- Can applying a counting semaphore **S** as a binary semaphore
- We can use semaphores to solve various synchronization problems

# Semaphore Usage Examples

## □ Example 1: Solution to the CS Problem

- Create a semaphore “**mutex**” initialized to 1

```
wait(mutex);
```

```
CS
```

```
signal(mutex);
```

## □ Example 2: Consider $P_1$ and $P_2$ that with two statements $S_1$ and $S_2$ and the requirement that $S_1$ to happen before $S_2$

- Create a semaphore “**synch**” initialized to 0

P1 :

```
 $S_1$ ;
```

```
signal(synch);
```

P2 :

```
wait(synch);
```

```
 $S_2$ ;
```

# Semaphore Implementation

- ❑ Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time
  - ❑ The implementation becomes the critical section problem where the **wait** and **signal** code are placed in the critical section
- ❑ Could now have **busy waiting** in critical section implementation
  - ❑ Little busy waiting if critical section rarely occupied
- ❑ Note that applications may spend lots of time in critical sections, so this is not a good solution

# Semaphore Implementation with **no Busy Waiting**

- Implementation of **wait**:

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        sleep();
    }
}
```

- Implementation of **signal**:

```
signal(semaphore *S) {
    S->value++;
    if (S->value <= 0 && S->list != NULL) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

- Semaphore structure

```
typedef struct {
    int value;
    struct process *list;
} semaphore;
```

# Synchronization Mechanisms: Monitors

# Monitors (管程)

- **Monitor:** A high-level abstraction that provides a convenient and effective mechanism for process synchronization
  - **Only one process may be active within the monitor at a time**
- Pseudocode syntax of a monitor:


```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }

    procedure P2 (...) { ... }

    procedure Pn (...) {.....}

    initialization code (...) { ... }
}
```

# Schematic view of a Monitor




# Condition Variables

- Monitors provide condition variables to allow threads to wait for specific conditions to become true before proceeding.
  - When a thread **waits** on a condition variable, **it releases the mutex**, allowing other threads to enter the monitor and potentially change the condition.
  - When the condition changes and another thread **signals** the condition variable, the waiting thread is **awakened and attempts to reacquire the mutex** before continuing its execution.
- Two operations are allowed on a condition variable:
  - **`x.wait()`** – a process that invokes the operation is suspended until **`x.signal()`**
  - **`x.signal()`** – resumes one of processes that invoked **`x.wait()`**
    - ▶ If no **`x.wait()`** on the variable, then it has no effect on the variable



# Monitor with Condition Variables



# Condition Variables Choices

- If process P invokes `x.signal()`, with Q in `x.wait()` state, what should happen next?
  - At most one process is allowed active within the monitor.
- Condition variable options:
  - **Signal and wait** – P waits until Q leaves monitor or waits for another condition
  - **Signal and continue** – Q waits until P leaves the monitor or waits for another condition
  - **Compromise** - P leaves the monitor immediately after executing signal, Q is resumed

# Implementation of Condition Variables

- For each condition variable  $x$ , we have:

```
semaphore x_sem; // condition variable semaphore
int x_count = 0; // #threads that are waiting for "x"
semaphore next; // hanging semaphore
int next_count = 0; // number of hanging processes
semaphore mutex; // for mutual exclusion
```

## Outer loop:

```
wait(mutex);
```

```
...
body of F
```

```
...
if (next_count > 0)
    signal(next);
else
    signal(mutex);
```

- The operation  $x.\text{wait}()$  can be implemented as:

```
x_count++;
if (next_count > 0)
    signal(next);
else
    signal(mutex);
wait(x_sem);
x_count--;
```

- The operation  $x.\text{signal}()$  can be implemented as:

```
if (x_count > 0) {
    next_count++;
    signal(x_sem);
    wait(next);
    next_count--;
}
```

# Monitor for Single-Resource Allocation

## **x.wait()** definition:

```
x_count++;  
if (next_count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x_sem);  
x_count--;
```

## **x.signal()** definition:

```
if (x_count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```

```
monitor ResourceAllocator  
{  
    boolean busy;  
    condition x;  
  
    void acquire(int time) {  
        if (busy)  
            x.wait(time);  
        busy = true;  
    }  
  
    void release() {  
        busy = false;  
        x.signal();  
    }  
    signal and wait  
  
    initialization_code() {  
        busy = false;  
    }  
}
```

# Monitor for Single-Resource Allocation

We will use the signal-and-wait scheme in our implementation. Since a signaling process must wait until the resumed process either leaves or waits, an additional binary semaphore, `next`, is introduced, initialized to 0. The signaling processes can use `next` to suspend themselves. An integer variable `next_count` is also provided to count the number of processes suspended on `next`. Thus, each external function `F` is replaced by

**管程的入口**

```
wait(mutex); 申请进入管程, 先在外部分列等待
...          成功进入, 执行自己的需求(函数F)
body of F    其中有可能涉及到资源x的wait和signal
...
管程的出口
if (next_count > 0) 执行完毕, 走之前抓一个别的
    signal(next);    进程来运行
else
    signal(mutex); 如果有人因signal-wait在等待,
                  优先恢复它; 否则, 放新人进来
```


Mutual exclusion within a monitor is ensured.

We can now describe how condition variables are implemented as well. For each condition `x`, we introduce a binary semaphore `x_sem` and an integer variable `x_count`, both initialized to 0. The operation `x.wait()` can now be implemented as

```
x.wait()
x_count++; 资源x的队列人数+1, 我马上进去了
if (next_count > 0)
    signal(next); 进去之前抓一个进程来运行, 同上
else
    signal(mutex);
wait(x_sem); 暂停, 进入资源x的等待队列
x_count--; 走到这说明有人把我唤醒了, 队列人数-1
```

The operation `x.signal()` can be implemented as

```
x.signal()
if (x_count > 0) { 是否有人在等待资源x? (由于 signal -wait)
    next_count++; 若有, 我将进入紧急等待队列
    signal(x_sem); 从x的资源队列中恢复一个进程
    wait(next); 暂停, 进入紧急等待队列
    next_count--; 有人把我唤醒了, 队列人数-1
}
```



# Resuming Processes within a Monitor

- If several processes queued on condition variable ***x***, and ***x.signal()*** is executed, which process should be resumed?
  - FCFS frequently not adequate
- Use the **conditional-wait** construct of the form

***x.wait(c)***

where:

- ***c*** is an integer (called the priority number)
- The process with lowest number (highest priority) is scheduled next

# Synchronization Problem Formulations

# Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
  - Bounded-Buffer Problem
  - Readers and Writers Problem
  - Dining-Philosophers Problem



# Improved Solution of Producer-Consumer Problem

```
while (true) {
```

**Producer**

```
    /* produce an item and put in nextProduced */
```

```
    while (counter == BUFFER_SIZE) ; // do nothing
```

```
    buffer [in] = nextProduced;
```

```
    in = (in + 1) % BUFFER_SIZE;
```

```
    wait(mutex);
```

```
    counter++;
```

```
    signal(mutex);
```

```
}
```

```
while (true) {
```

**Consumer**

```
    while (counter == 0) ; // do nothing
```

```
    nextConsumed = buffer[out];
```

```
    out = (out + 1) % BUFFER_SIZE;
```

```
    wait(mutex);
```

```
    counter--;
```

```
    signal(mutex);
```

```
    /* consume the item */
```

```
}
```

# Sync Problems: (1) Bounded-Buffer Problem

- $n$  buffers, each can hold one item
- Semaphore `mutex` initialized to the value 1: Control access to the shared buffer.
- Semaphore `full` initialized to the value 0: Produced item number.
- Semaphore `empty` initialized to the value  $n$ : Idle slot number.

## □ Producer process

```
while (true) {  
    /* produce an item in  
    next_produced */  
    wait(empty);  
    wait(mutex);  
    /* add next produced to  
    the buffer */  
    signal(mutex);  
    signal(full);  
}
```

## □ Consumer process

```
while (true) {  
    wait(full);  
    wait(mutex);  
    /* remove an item from buffer  
    to next_consumed */  
    signal(mutex);  
    signal(empty);  
    /* consume the item in next  
    consumed */  
}
```

# Sync Problems: (2) Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - **Readers** – only read the data set; **they do not perform any updates**
  - **Writers** – can both read and write
- Problem defined:
  - **Allow multiple readers to read at the same time**
  - Only one single writer can access the shared data at each point of time
- Shared Data

| Type      | Name              | Init. Value | Role                                                                                   |
|-----------|-------------------|-------------|----------------------------------------------------------------------------------------|
| Semaphore | <b>rw_mutex</b>   | 1           | Ensure mutual modification to the data set and mutual-exclusion of reading and writing |
| INT       | <b>read_count</b> | 0           | Record # of concurrent readers                                                         |
| Semaphore | <b>mutex</b>      | 1           | Ensure mutual access to read_count                                                     |

# Sync Problems: (2) Readers-Writers Problem

## □ Writer process

```
while (true) {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
}
```

Goal: Allow multiple readers to read at the same time, but one writer each time.

## □ Reader process

```
while (true){  
    wait(mutex);  
    read_count++;  
    if (read_count == 1) /* first reader */  
        wait(rw_mutex);  
    signal(mutex);  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0) /* last reader */  
        signal(rw_mutex);  
    signal(mutex);  
}
```

# Sync Problems: (2) Readers-Writers Problem Variations

- ❑ The solution in previous slide can result in a starvation situation where a writer process may never writes.
  - ❑ It is a reader-preferred solution
- ❑ **Reader-Preferred Solution** – no reader kept waiting unless writer has permission to use shared object
  - ❑ No reader should wait for other readers to finish simply because a writer is waiting
- ❑ **Writer-Preferred Solution** – once writer is ready, it performs write asap
  - ❑ If a writer is waiting to access the object, no new readers may start reading

# Sync Problems: (2) Writer-Preferred Solution

```
int read_count = 0, write_count = 0;
semaphore mutexrc = 1, mutexwc = 1, wrt = 1, rd = 1;
```

□ Writer process

```
do {
    wait(mutexwc);
    writecount++;
    if (writecount == 1)
        wait(rd);
    signal(mutexwc);

    wait(wrt);
    // writing is performed
    signal(wrt);

    wait(mutexwc);
    writecount--;
    if (writecount == 0)
        signal(rd);
    signal(mutexwc);
} while(TRUE);
```

□ Reader process

```
do {
    wait(rd);
    wait(mutexrc);
    readcount++;
    if (readcount == 1)
        wait(wrt);
    signal(mutexrc);
    signal(rd);

    //reading is performed

    wait(mutexrc);
    readcount--;
    if (readcount == 0)
        signal(wrt);
    signal(mutexrc);
} while(TRUE);
```

# Readers-Writers Problem Variations

- *Reader-Preferred Solution* – no reader kept waiting unless writer has permission to use shared object
  - no reader should wait for other readers to finish simply because a writer is waiting
- *Writer-Preferred Solution* – once writer is ready, it performs write asap
  - if a writer is waiting to access the object, no new readers may start reading
- Both may have starvation leading to even more variations
- Find a solution to starvation-free reader-writer problem!

# No-Starvation Solution

```
int readcount = 0;
semaphore mutex = 1, wrt = 1, rd = 1;
```

□ Writer process

```
do {
    wait ( wrt );
    wait ( rd );

    // writing is performed

    signal ( rd );
    signal ( wrt );
} while (TRUE);
```

□ Reader process

```
do {
    wait ( wrt );
    wait ( mutex );
    readcount ++;
    if (readcount == 1)
        wait (rd);
    signal ( mutex );
    signal ( wrt );


    //reading is performed

    wait ( mutex );
    readcount --;
    if (readcount == 0)
        signal (rd);
    signal ( mutex );
} while (TRUE);
```



# Sync Problems: (3) Dining-Philosophers Problem

- N philosophers (哲学家) sit at a round table with a bowl of rice in the middle.
  - They spend their lives thinking and eating
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
  - Need both to eat, then release both when done
- In the case of 5 philosophers
  - Shared data
    - ▶ Chopsticks
    - ▶ Semaphore **chopstick** [5] initialized to 1



# Sync Problems: (3) Dining-Philosophers with Semaphores

- Semaphore Solution
- The structure of Philosopher  $i$  :

```
while(true) {  
    wait(chopstick[i] );  
    wait(chopstick[(i + 1)%5]);  
  
    /* eat for awhile */  
  
    signal(chopstick[i] );  
    signal(chopstick[(i + 1)%5]);  
  
    /* think for awhile */  
  
}
```

- What is the problem with this algorithm?
  - Deadlocks!

# Sync Problems: (3) Dining-Philosophers with Monitors

- Deadlock-free solution with monitors.
- Three states for each philosopher:
  - `enum {THINKING, HUNGRY, EATING}`  
`state[5];`
- Condition variables for all philosophers:
  - `condition self[5];`
- Eat condition for a philosopher:
  - Two neighbors are not eating  
`(state[(i+4) % 5] != EATING) &`  
`(state[(i+1) % 5] != EATING)`
- Limitation:
  - One philosopher may stare to death

```
monitor DiningPhilosophers
{
    enum {THINKING, HUNGRY, EATING} state[5];
    condition self[5];

    void pickup(int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait();
    }

    void putdown(int i) {
        state[i] = THINKING;
        test((i + 4) % 5);
        test((i + 1) % 5);
    }

    void test(int i) {
        if ((state[(i + 4) % 5] != EATING) &&
            (state[i] == HUNGRY) &&
            (state[(i + 1) % 5] != EATING)) {
            state[i] = EATING;
            self[i].signal();
        }
    }

    initialization_code() {
        for (int i = 0; i < 5; i++)
            state[i] = THINKING;
    }
}
```

Acquire

Release

# Deadlock and Starvation

- ❑ **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the **waiting processes**
- ❑ Let **S** and **Q** be two semaphores initialized to 1

$P_0$

① wait (S);

③ wait (Q);

.

.

signal (S);

signal (Q);

$P_1$

② wait (Q);

④ wait (S);

.

.

signal (Q);

signal (S);

- ❑ **Starvation** – indefinite blocking
  - ❑ A process may never be removed from the semaphore queue in which it is suspended

# Summary

- ❑ A critical section is a code section where shared data may be manipulated and a possible race condition may occur.
- ❑ A solution to the critical-section problem must satisfy: (1) mutual exclusion, (2) progress, and (3) bounded waiting.
- ❑ Peterson's solution does not work well on modern computer architectures.
- ❑ Hardware solutions to the critical-section problems: (1) memory barriers, (2) hardware instruction (e.g., compare-and-swap), and (3) atomic variables.
- ❑ Mutex and semaphores can be used to provide mutual exclusion
- ❑ A monitor uses condition variables to allow processes to wait for certain conditions and to signal one another when conditions are true.
- ❑ Classic problems of process synchronization include the bounded-buffer, readers–writers, and dining-philosophers problems.
  - ❑ Solutions can be developed using mutex locks, semaphores, monitors, and condition variables.

# Homework

- Reading
  - Chapter 6
  - Chapter 7
- HW1 due on Mar 12 at 23:59!