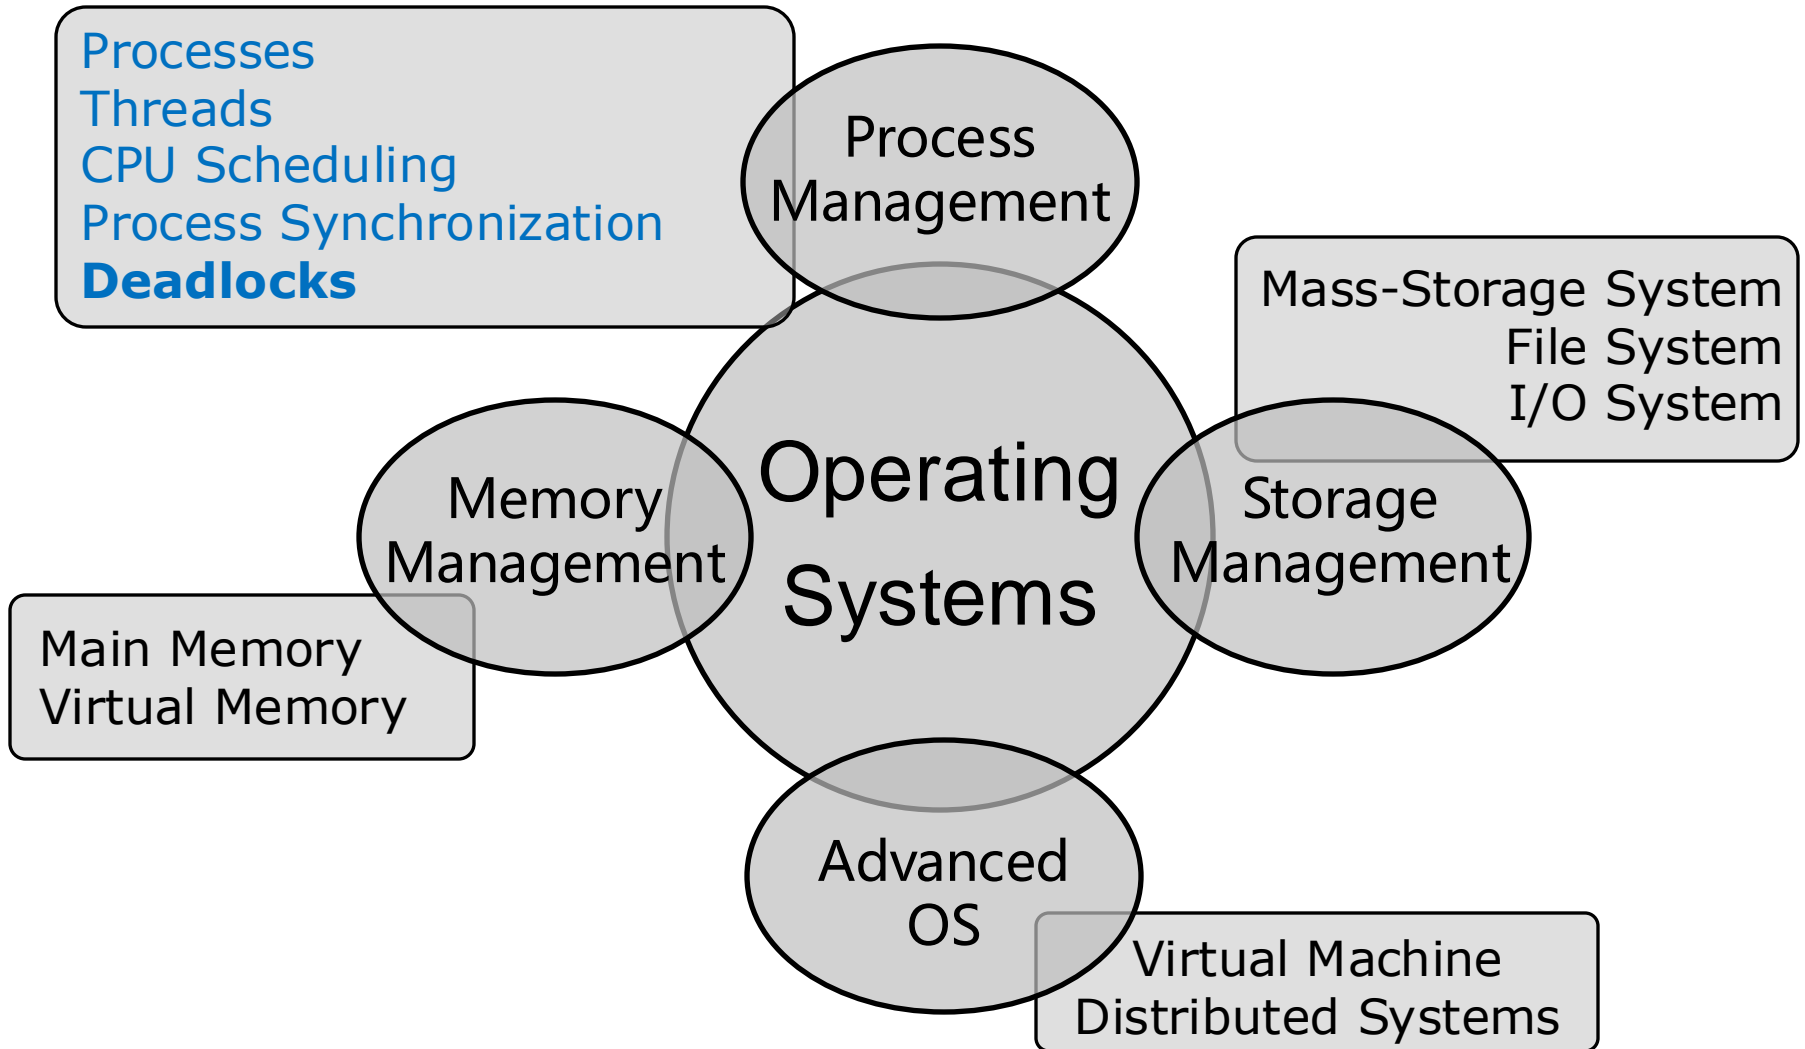


# Deadlocks

**Shengzhong Liu**

Department of Computer Science and Engineering  
Shanghai Jiao Tong University

# Operating System Topics




# Outline

- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
  - Never enters a deadlock:
    - ▶ Deadlock Prevention
    - ▶ Deadlock Avoidance
  - Enters and recovers from the deadlock:
    - ▶ Deadlock Detection
    - ▶ Recovery from Deadlock

# System Model

# Analogy: Bridge Crossing Example



- Traffic only in one direction
- Each section of a bridge can be viewed as a resource
- A **deadlock** occurs when two cars get on the bridge from different directions at the same time

# Deadlock with Semaphores

- Data:
  - A semaphore  $s_1$  initialized to 1
  - A semaphore  $s_2$  initialized to 1
- Assume we have two threads  $T_1$  and  $T_2$
- |   |   |
|---|---|
| □ $T_1$ :<br><br><code>wait(s<sub>1</sub>)</code><br><br><code>wait(s<sub>2</sub>)</code> | □ $T_2$ :<br><br><code>wait(s<sub>2</sub>)</code><br><br><code>wait(s<sub>1</sub>)</code> |
|---|---|
- **Deadlock**: A set of blocked threads each holding some resources and waiting to acquire the resources held by another thread in the set

# System Model

## □ Notations:

- Threads  $T_1, T_2, \dots, T_n$

- Resource types  $R_1, R_2, \dots, R_m$

e.g., CPU, memory space, I/O devices, mutex and semaphores

- **Each resource type  $R_i$  has  $W_i$  instances.**

## □ Each process utilizes a resource in the following sequence:

- **Request:** Wait if the request cannot be granted immediately.

- **Use:** The process operates on the resource.

- **Release:** The process releases the resource.

## □ Deadlock State:

- A set of threads is in a deadlocked state when every thread in the set is waiting for an event that can be caused only by another thread in the set.

- Events mainly include resource acquisition and release.


# Deadlock Characterization



# Deadlock Characterization


Deadlock can arise if four conditions hold simultaneously.

- ❑ **Mutual exclusion:** only one thread at a time can use a resource
- ❑ **Hold and wait:** a thread holding at least one resource is waiting to acquire additional resources held by other threads
- ❑ **No preemption:** a resource can be only released voluntarily by the thread after its task has completed
- ❑ **Circular wait:** there exists a set  $\{T_0, T_1, \dots, T_n\}$  of waiting threads such that
  - ❑  $T_0$  is waiting for a resource held by  $T_1$
  - ❑  $T_1$  is waiting for a resource held by  $T_2$
  - ❑ ...
  - ❑  $T_{n-1}$  is waiting for a resource held by  $T_n$
  - ❑  $T_n$  is waiting for a resource held by  $T_0$




# Resource-Allocation Graph

- Deadlocks can be identified with **system resource-allocation graph**.
  - A set of vertices **V** and a set of edges **E**.
- V is partitioned into two types:
  - $T = \{T_1, T_2, \dots, T_n\}$  is the set consisting of all the threads in the system.
  - $R = \{R_1, R_2, \dots, R_m\}$  is the set consisting of all resource types in the system
- There are two types of edges:
  - Request edge** – directed edge  $T_i \rightarrow R_j$
  - Assignment edge** – directed edge  $R_j \rightarrow T_i$



# Example: Resource Allocation Graph

- Threads:  $T = \{T_1, T_2, T_3\}$
- Resources:  $R = \{R_1, R_2, R_3, R_4\}$ 
  - Resource instances:
    - ▶  $R_1$  x1,  $R_2$  x2,  $R_3$  x1,  $R_4$  x3
- Edges:  $E = \{T_1 \rightarrow R_1, T_2 \rightarrow R_3, R_1 \rightarrow T_2, R_2 \rightarrow T_2, R_2 \rightarrow T_1, R_3 \rightarrow T_3\}$ 
  - $T_1$ :
    - ▶ holds one instance of  $R_2$
    - ▶ is waiting for an instance of  $R_1$
  - $T_2$ :
    - ▶ holds one instance of  $R_1$ , one instance of  $R_2$
    - ▶ is waiting for an instance of  $R_3$
  - $T_3$ :
    - ▶ holds one instance of  $R_3$




# Basic Facts

- If graph **contains no circle**  $\Rightarrow$  No deadlock
- If graph **contains a circle**  $\Rightarrow$ 
  - If only one instance per resource type, then deadlock
  - If several instances per resource type, possibility of deadlock
    - ▶ Cycle is a necessary but not sufficient condition for deadlocks.
- Question:
  - Can you find a way to determine whether there is a deadlock, given a resource allocation graph with several instances per resource type?

# Example: Resource Allocation Graph w/ Deadlock(s)

□ One circle


□  $T_1 \rightarrow R_1 \rightarrow T_2 \rightarrow R_3 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1$




□ Two circles

□  $T_1 \rightarrow R_1 \rightarrow T_2 \rightarrow R_3 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1$

□  $T_2 \rightarrow R_3 \rightarrow T_3 \rightarrow R_2 \rightarrow T_2$



# Example: Graph with A Cycle w/o Deadlock



# Methods for Handling Deadlocks

# Methods for Handling Deadlocks

- **Option 1:** Ensure that the system will **never** enter a deadlock state
  - **Deadlock prevention**
    - ▶ 【Pessimistic policy, no prior information】
    - ▶ Ensure that at least one necessary condition cannot hold.
    - ▶ Conservative resource allocation, but may lead to underutilization
  - **Deadlock avoidance**
    - ▶ 【Dynamic estimation, use prior information】
    - ▶ Employ strategies to ensure the system does not enter deadlock at each step of resource allocation
- **Option 2:** **Allow** the system to enter a deadlock state and then recover
  - **Deadlock detection**
  - **Deadlock recovery**



# Methods for Handling Deadlocks: Deadlock Prevention

# Deadlock Prevention

Examine the **necessary conditions for deadlocks** to happen:

- **Mutual Exclusion** – removed for sharable resources (e.g., read-only files); must hold for non-sharable resources
  - We can not prevent deadlocks by denying the mutual-exclusion condition.
- **Hold and Wait** – when a thread requests a resource, it can not hold any other resources
  - Option 1: Require a thread to request and be allocated all its resources before it begins execution
    - ▶ Impractical for most applications.
  - Option 2: Allow a thread to request resources only when the thread has none allocated to it.
    - ▶ Before requesting additional resources, a thread must release all allocated resources.
  - Limitations: Low resource utilization; starvation possible

# Deadlock Prevention (Cont.)

## □ No Preemption:

- Support resource preemption:
  - ▶ If a thread is holding some resources and requesting another resource that can not be allocated, then release all its held resources
  - ▶ Preempted resources are added to the list of waiting resources
  - ▶ The thread is restarted only when it can regain its old resources and new resources
- Often applied to resources whose state can be easily saved and restored later, such as CPU registers and database transactions.
  - ▶ Can not be applied to mutex locks and semaphores.

## □ Circular Wait:

- Impose a total ordering of all resource types, and require that each thread requests resources in an increasing order of enumeration
- The only practical option.

# Break Circular Wait with Lock Order

## □ Policy:

- Invalidating the circular wait condition is most common.
- Simply assign each resource (i.e., mutex locks) a unique number.
- Resources must be acquired in order.

## □ If:

`first_mutex = 1`  
`second_mutex = 5`

code for `thread_two` could not be written as follows:



```
/* thread_one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}
```

```
-----
/* thread_two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```

# Methods for Handling Deadlocks: Deadlock Avoidance

# Deadlock Avoidance


Requires the system to have some additional **prior information**

- Simplest model: Each thread declare the **maximum #resources of each needed type**
- **Avoidance**: dynamically check the resource-allocation state, ensure there can never be a circular-wait condition
- **Resource-allocation state**:
  - Number of available resources
  - Number of allocated resources
  - Maximum demands of the threads

**Additional knowledge required!**


# Safe State

- When a thread requests an available resource, must decide if the allocation leaves the system in a safe state
- System is in **safe state** if there exists a sequence  $\langle T_1, T_2, \dots, T_n \rangle$  of ALL the threads in the systems such that
  - For each  $T_i$ , the resources that it will request can be satisfied by currently available resources + resources held by all the  $T_j$ , with  $j < i$
- That is:
  - If  $T_i$  resource needs are not available,  $T_i$  waits until all  $T_j$  have finished
  - When  $T_j$  is finished,  $T_i$  can obtain its resources, execute, return all resources, and terminate
  - When  $T_i$  terminates,  $T_{i+1}$  can obtain its needed resources, and so on



# Safe, Unsafe, Deadlock State

- If a system is in safe state  
⇒ no deadlocks
- If a system is in unsafe state  
⇒ possibility of deadlock
- Avoidance  
⇒ ensure that a system will never enter an unsafe state.





# Safe & Unsafe States

	Maximum Needs	Holds	Needs
$T_0$	10	5	5
$T_1$	4	2	2
$T_2$	9	2	7

Available
3

Safe sequence: ?

# Safe & Unsafe States

	Maximum Needs	Holds	Needs
$T_0$	10	5	5
$T_1$	4	4	0
$T_2$	9	2	7



Available
1

Safe sequence:  $T_1$

# Safe & Unsafe States

	Maximum Needs	Holds	Needs
$T_0$	10	5	5
$T_1$	4	--	--
$T_2$	9	2	7

Available
5

Safe sequence:  $T_1$

# Safe & Unsafe States

	Maximum Needs	Holds	Needs
T <sub>0</sub>	10	10	0
T <sub>1</sub>	4	--	--
T <sub>2</sub>	9	2	7



Available
0

Safe sequence: T<sub>1</sub> → T<sub>0</sub>

# Safe & Unsafe States

	Maximum Needs	Holds	Needs
$T_0$	10	--	--
$T_1$	4	--	--
$T_2$	9	2	7

Available
10

Safe sequence:  $T_1 \rightarrow T_0$

# Safe & Unsafe States

	Maximum Needs	Holds	Needs
T <sub>0</sub>	10	--	--
T <sub>1</sub>	4	--	--
T <sub>2</sub>	9	9	0



Available
3

Safe sequence: T<sub>1</sub> → T<sub>0</sub> → T<sub>2</sub>

# Safe & Unsafe States

	Maximum Needs	Holds	Needs
$T_0$	10	--	--
$T_1$	4	--	--
$T_2$	9	--	--

Available
12

Safe sequence:  $T_1 \rightarrow T_0 \rightarrow T_2$

# Safe & Unsafe States

Back from beginning.

	Maximum Needs	Holds	Needs
$T_0$	10	5	5
$T_1$	4	2	2
$T_2$	9	2	7

Available
3

Safe sequence: ?



# Safe & Unsafe States

Give T2 one more?

	Maximum Needs	Holds	Needs
T <sub>0</sub>	10	5	5
T <sub>1</sub>	4	2	2
T <sub>2</sub>	9	3	6

**Available**

2

Safe sequence: ?

# Safe & Unsafe States

	Maximum Needs	Holds	Needs
$T_0$	10	5	5
$T_1$	4	--	--
$T_2$	9	3	6

Available
4


Safe sequence:  $T_1 \rightarrow ?$

# Deadlock Avoidance Algorithms

- Avoidance algorithms ensure the system will never deadlock
  - When a process requests a resource, the request is granted only if **the allocation leaves the system in a safe state**.
- Two avoidance algorithms
  - If we have a single instance of a resource type
    - ▶ Use a **resource-allocation-graph algorithm**
  - If we have multiple instances of a resource type
    - ▶ Use the **banker's algorithm**


# Resource-Allocation-Graph Algorithm

- **[New]: Claim edge**  $T_i \rightarrow R_j$  indicates that thread  $T_i$  **may** request resource  $R_j$ ; represented by a **directed dashed line**
  - **Claim edge** converts to **request edge** when a thread requests a resource
  - **Request edge** converts to an **assignment edge** when the resource is allocated to the thread
- When a resource is released by a thread, assignment edge reconverts to a claim edge (the edge is removed if the thread finishes)




# Resource-Allocation-Graph Algorithm

- Suppose that thread  $T_i$  requests a resource  $R_j$
- The request can be granted only if
  - Converting the request edge to an assignment edge does not result in a circle in the resource allocation graph



Can we grant  $T_2$ 's request for  $R_2$ ?



Circle! Therefore,  $T_2$ 's request cannot be granted, and  $T_2$  needs to wait.

# Banker's Algorithm

- We have multiple instances of each resource
- Each thread must claim its maximum use of each resource in advance
- When a thread requests a resource, it may have to wait
- When a process gets all its resources it must return them in a finite amount of time

# Data Structures for the Banker's Algorithm

Let  $n$  = number of threads, and  $m$  = number of resource types.

- **Available:  $m$  vector.** The #available resources of each type.
  - If  $\text{available}[j] = k \rightarrow k$  instances of resource type  $R_j$  are available
- **Max:  $n \times m$  matrix.** Maximum demand of each thread.
  - If  $\text{Max}[i, j] = k \rightarrow$  thread  $T_i$  may request at most  $k$  of  $R_j$
- **Allocation:  $n \times m$  matrix.** #allocated resource of each type.
  - If  $\text{Allocation}[i, j] = k \rightarrow T_i$  is currently allocated  $k$  of  $R_j$
- **Need:  $n \times m$  matrix.** The remaining resource need of each thread.
  - If  $\text{Need}[i, j] = k \rightarrow P_i$  may need  $k$  more  $R_j$  to complete
  - $\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$

# Banker's Algorithm

- **【Variable】 Work:** The available resources can be allocated to processes at any given point during execution.
  - As the algorithm progresses, the "work" variable is used to simulate resource allocation to threads.
- **【Status】 Finish:** A boolean array that indicates whether each thread can complete its execution without leading to deadlock.
  - As the algorithm proceeds, the "finish" array is updated to reflect which processes can potentially complete execution safely.
- **Key idea:**
  - Employs a series of **checks and simulations** to ensure resource allocation will not result in a deadlock.
  - If **a safe sequence exists**, resources can be allocated to threads without risking deadlock.
  - **Otherwise**, the system should wait until resources become available or deny the request to avoid potential deadlock.



# Subroutine: Safety Check Algorithm

Find out whether or not the system is in a safe state.

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively.
  - **Work = Available**
  - **Finish[i] = false**, for  $i = 0, 1, \dots, n-1$
2. Find an  $i$  such that both:
  - (a) **Finish[i] = false**
  - (b) **Need<sub>i</sub> ≤ Work**If no such  $i$  exists, go to step 4
3. **Work = Work + Allocation<sub>i</sub>**  
**Finish[i] = true**  
go to step 2
4. If **Finish[i] == true for all i**, then the system is in a safe state;  
otherwise, the system is in a deadlock

# Resource-Request Algorithm for Thread $T_i$

Determine whether the requests can be safely granted.

**Request<sub>i</sub>** = request vector for thread  $T_i$ . If **Request<sub>i</sub> [j] = k** then thread  $T_i$  wants  $k$  instances of resource type  $R_j$

1. If **Request<sub>i</sub> ≤ Need<sub>i</sub>**, go to step 2. Otherwise, raise error condition, since thread has exceeded its maximum claim
2. If **Request<sub>i</sub> ≤ Available**, go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

**Available = Available – Request<sub>i</sub>;**

**Allocation<sub>i</sub> = Allocation<sub>i</sub> + Request<sub>i</sub>;**

**Need<sub>i</sub> = Need<sub>i</sub> – Request<sub>i</sub>;**

- Test the safety of the new system state:
  - If **safe**  $\Rightarrow$  the resources are allocated to  $P_i$
  - If **unsafe**  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

# Example of Banker's Algorithm

□ **5 threads**  $T_0$  through  $T_4$ ;

**3 resource types:**

A (**10** instances), B (**5** instances), and C (**7** instances)

Initial system snapshot:

	Max	Allocation	Need	Available
	A B C	A B C	A B C	A B C
$T_0$	7 5 3	0 1 0	7 4 3	3 3 2
$T_1$	3 2 2	2 0 0	1 2 2	
$T_2$	9 0 2	3 0 2	6 0 0	
$T_3$	2 2 2	2 1 1	0 1 1	
$T_4$	4 3 3	0 0 2	4 3 1	

□ **Is the system in safe state?**

# Applying Safety Algorithm

	Max	Allocation	Need	Available
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
$T_0$	7 5 3	0 1 0	7 4 3	5 3 2
$T_2$	9 0 2	3 0 2	6 0 0	
$T_3$	2 2 2	2 1 1	0 1 1	
$T_4$	4 3 3	0 0 2	4 3 1	



Safe sequence:  $T_1$

# Applying Safety Algorithm

	Max	Allocation	Need	Available
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
$T_0$	7 5 3	0 1 0	7 4 3	7 4 3
$T_2$	9 0 2	3 0 2	6 0 0	
$T_4$	4 3 3	0 0 2	4 3 1	



Safe sequence:  $T_1 \rightarrow T_3$

# Applying Safety Algorithm

	Max	Allocation	Need	Available
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
				7 5 3
$T_2$	9 0 2	3 0 2	6 0 0	
$T_4$	4 3 3	0 0 2	4 3 1	



Safe sequence:  $T_1 \rightarrow T_3 \rightarrow T_0$

# Applying Safety Algorithm

	Max	Allocation	Need	Available
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
				10 5 5
$T_4$	4 3 3	0 0 2	4 3 1	



Safe sequence:  $T_1 \rightarrow T_3 \rightarrow T_0 \rightarrow T_2$

# Applying Safety Algorithm

	Max	Allocation	Need	Available
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
				10 5 7



Safe sequence:  $T_1 \rightarrow T_3 \rightarrow T_0 \rightarrow T_2 \rightarrow T_4$

Safe!



# Example: $T_1$ Request (1,0,2)

- Check that Request  $\leq$  Available (that is,  $T_1 \rightarrow (1,0,2) \leq (3,3,2) \Rightarrow \text{true}$ )

	Max	Allocation	Need	Available
	A B C	A B C	A B C	A B C
$T_0$	7 5 3	0 1 0	7 4 3	2 3 0
$T_1$	3 2 2	3 0 2	0 2 0	
$T_2$	9 0 2	3 0 2	6 0 0	
$T_3$	2 2 2	2 1 1	0 1 1	
$T_4$	4 3 3	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence  $\langle T_1, T_3, T_0, T_2, T_4 \rangle$  satisfies safety requirement
- Grant the request by  $T_1$

# Example: $T_0$ Request (0,2,0)

- Check that Request  $\leq$  Available (that is,  $T_0 \rightarrow (0,2,0) \leq (2,3,0) \Rightarrow \text{true}$ )

	Max	Allocation	Need	Available
	A B C	A B C	A B C	A B C
$T_0$	7 5 3	0 3 0	7 2 3	2 1 0
$T_1$	3 2 2	3 0 2	0 2 0	
$T_2$	9 0 2	3 0 2	6 0 0	
$T_3$	2 2 2	2 1 1	0 1 1	
$T_4$	4 3 3	0 0 2	4 3 1	

- Does there exist a safe sequence exist?
  - **No!**
- The request should be held and wait.


# Methods for Handling Deadlocks: Deadlock Detection

# Deadlock Detection

- Allow the system to enter a deadlock state
- Deadlock detection algorithm
- Deadlock recovery scheme


# Single Instance of Each Resource Type

- Maintain *wait-for* graph
  - Nodes are threads
  - $T_i \rightarrow T_j$  if  $T_i$  is waiting for  $T_j$
- Periodically invoke an algorithm that searches for a circle in the graph. If there is a circle, there exists a deadlock



(a)

Resource-Allocation Graph



(b)

Corresponding wait-for graph

# Several Instances of a Resource Type

- **Available:** A vector of length  $m$  indicates the number of available resources of each type.
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each thread.
- **Request:** An  $n \times m$  matrix indicates the current request of each thread. If  $\text{Request}[i][j] = k$ , then thread  $T_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Detection Algorithm

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , and initialize:
  - (a) **Work** = Available
  - (b) For  $i = 1, 2, \dots, n$ , if  $\text{Allocation}_i \neq 0$ , then  $\text{Finish}[i] = \text{false}$ ; otherwise,  $\text{Finish}[i] = \text{true}$
2. Find an index  $i$  such that both:

- (a)  **$\text{Finish}[i] == \text{false}$**
  - (b)  **$\text{Request}_i \leq \text{Work}$**

If no such  $i$  exists, go to step 4
3.  **$\text{Work} = \text{Work} + \text{Allocation}_i$**   
 **$\text{Finish}[i] = \text{true}$**   
go to step 2
4. If  $\text{Finish}[i] == \text{false}$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, **if  $\text{Finish}[i] == \text{false}$ , then  $P_i$  is deadlocked**

**Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state**

# Example of Detection Algorithm

- Five threads  $T_0$  through  $T_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)

- Initial snapshot:

	Allocation	Request	Available
	A B C	A B C	A B C
$T_0$	0 1 0	0 0 0	0 0 0
$T_1$	2 0 0	2 0 2	
$T_2$	3 0 3	0 0 0	
$T_3$	2 1 1	1 0 0	
$T_4$	0 0 2	0 0 2	

- Sequence  $\langle T_0, T_2, T_3, T_1, T_4 \rangle$  will result in  $\text{Finish}[i] = \text{true}$  for all  $i$



## Example (Cont.)

- $T_2$  requests an additional instance of type C

	Allocation	Request	Available
	A B C	A B C	A B C
$T_0$	0 1 0	0 0 0	0 0 0
$T_1$	2 0 0	2 0 2	
$T_2$	3 0 3	0 0 1	
$T_3$	2 1 1	1 0 0	
$T_4$	0 0 2	0 0 2	

- State of the system?
  - Can reclaim resources held by thread  $T_0$ , but insufficient resources to fulfill other threads' requests
  - Deadlock exists, consisting of threads  $T_1$ ,  $T_2$ ,  $T_3$ , and  $T_4$

# Detection-Algorithm Usage

- When and how often to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many threads will need to be rolled back?
- If the detection algorithm is invoked for every resource request, considerable overhead in computation time will be incurred.
- If the detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph
  - We would not be able to tell which of the many deadlocked threads “caused” the deadlock.

# Methods for Handling Deadlocks: Recovery from Deadlocks

# Recovery from Deadlock

- Process/Thread **Termination**

- Abort one or more processes to break the circular wait

- Resource **Preemption**

- Preempt some resources from one or more of the deadlocked threads

# Process Termination

- ❑ Abort all deadlocked processes
- ❑ Abort one process at a time until the deadlock cycle is eliminated
- ❑ In which order should we choose to abort?
  - ❑ Priority of the process
  - ❑ How long process has computed, and how much longer to completion
  - ❑ Resources the process has used
  - ❑ Resources process needs to complete
  - ❑ How many processes will need to be terminated
  - ❑ Is process interactive or batch?

# Resource Preemption

- **Selecting a victim** – minimize cost
  - We must determine the order of preemption to minimize cost
- **Rollback** – return to some safe state, restart process from that state
- **Starvation** – same process may always be picked as victim, include number of rollback in cost factor
  - We must ensure that a process can be picked as a victim only a (small) finite number of times
  - We can include number of rollbacks in the cost factor

# Summary

- ❑ Deadlock occurs in a set of threads when every thread in the set is waiting for an event that can only be caused by another thread in the set
- ❑ Four necessary conditions for deadlock: (1) mutual exclusion, (2) hold and wait, (3) no preemption, and (4) circular wait.
  - ❑ Deadlock is only possible when all four conditions are present
- ❑ Deadlocks can be modeled with **resource-allocation graphs**, where a cycle indicates deadlock.
- ❑ **Deadlock prevention**: Ensuring one of the four necessary conditions cannot occur
- ❑ **Deadlock avoidance**: Evaluate threads and resources to determine if the system is in a deadlocked state
- ❑ **Deadlock recovery**: Process termination or resource preemption.

# Homework

- Reading
  - Chapter 8
- Check Canvas for HW2 release, due on **Mar. 21 at 23:59!**