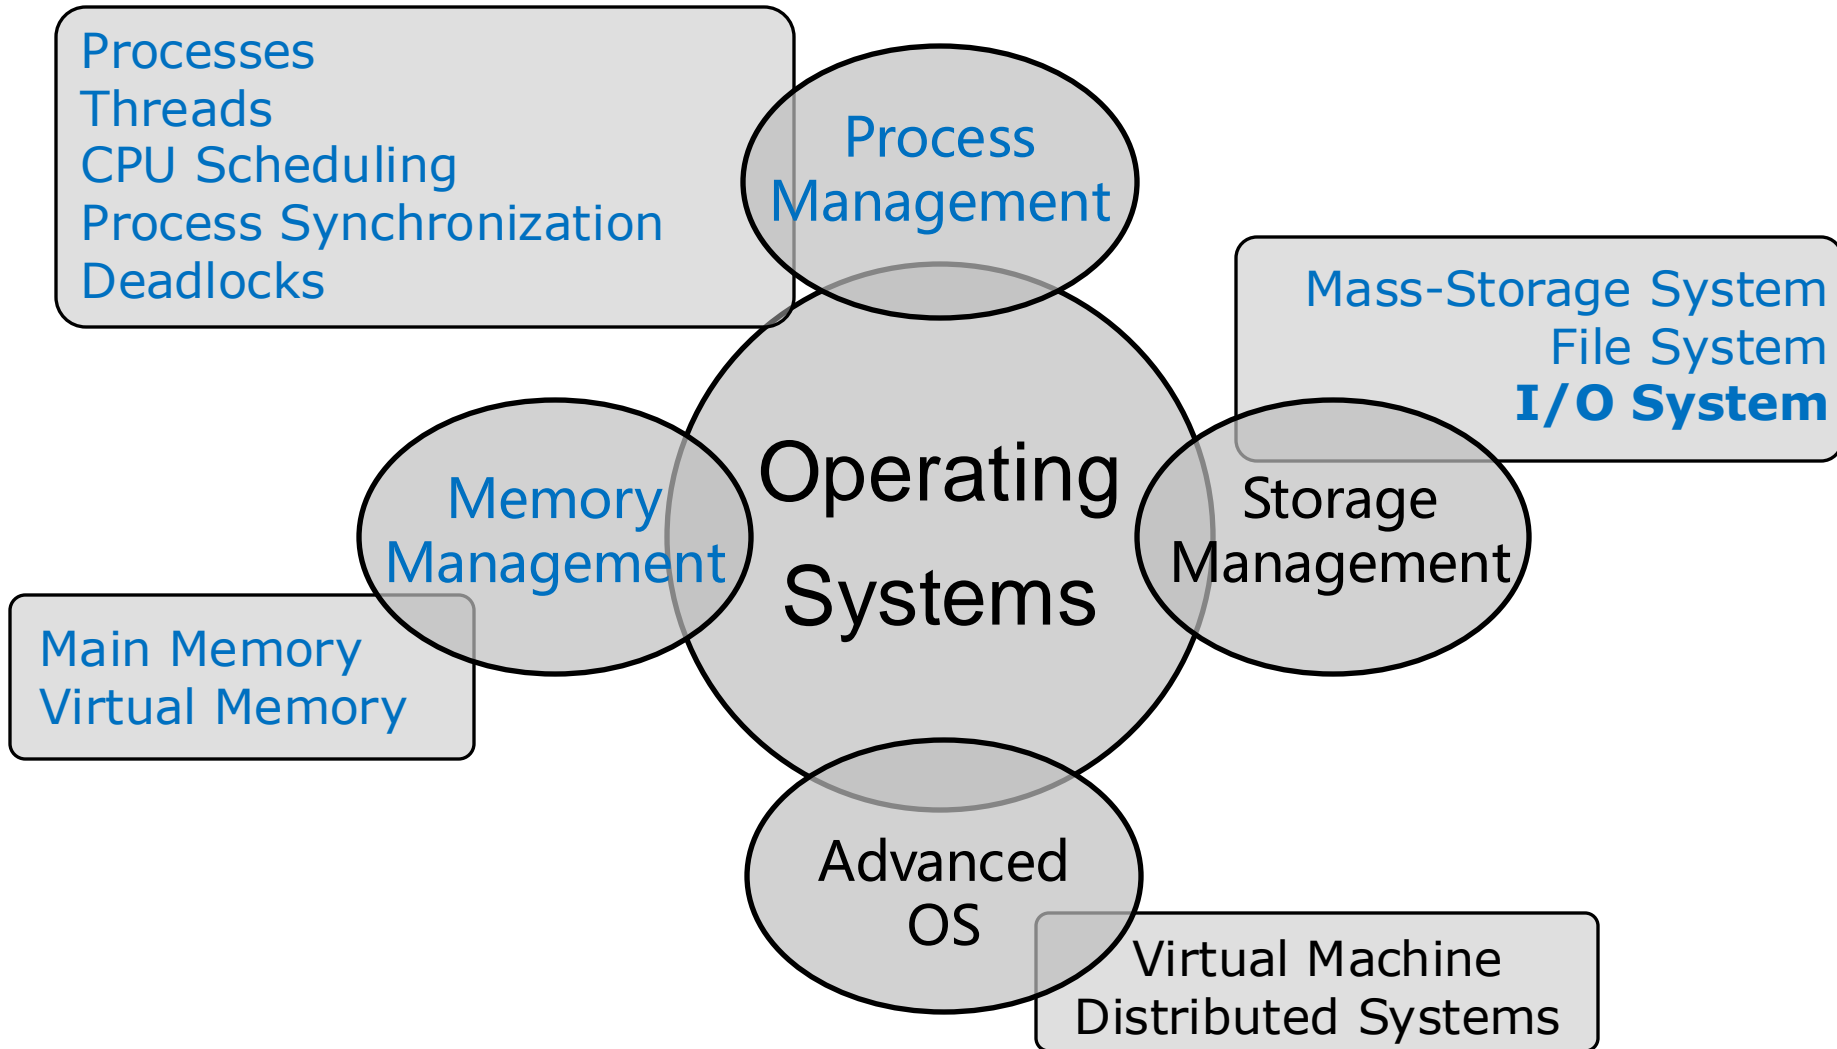


I/O Systems

Shengzhong Liu

Department of Computer Science and Engineering
Shanghai Jiao Tong University

Operating System Topics



Outline

- Overview
- I/O Hardware
- Application I/O Interface
- Kernel I/O Subsystem
- Transforming I/O Requests to Hardware Operations

Overview

- I/O management is a major component of OS design and operation
 - I/O devices vary greatly, and new types of devices frequently appear
 - Various methods to control them
 - Performance management
- **Ports, buses, and device controllers** connect to various devices
- **Device drivers (设备驱动)** encapsulate device details
 - Present uniform device-access interface to I/O subsystem

I/O Hardware

I/O Device Types




Keyboard/Trackpad/Mouse



Touchscreen



Scanner



Microphone



Webcam



Monitor



Printer



Projector



Headphone



Plotter



Joystick

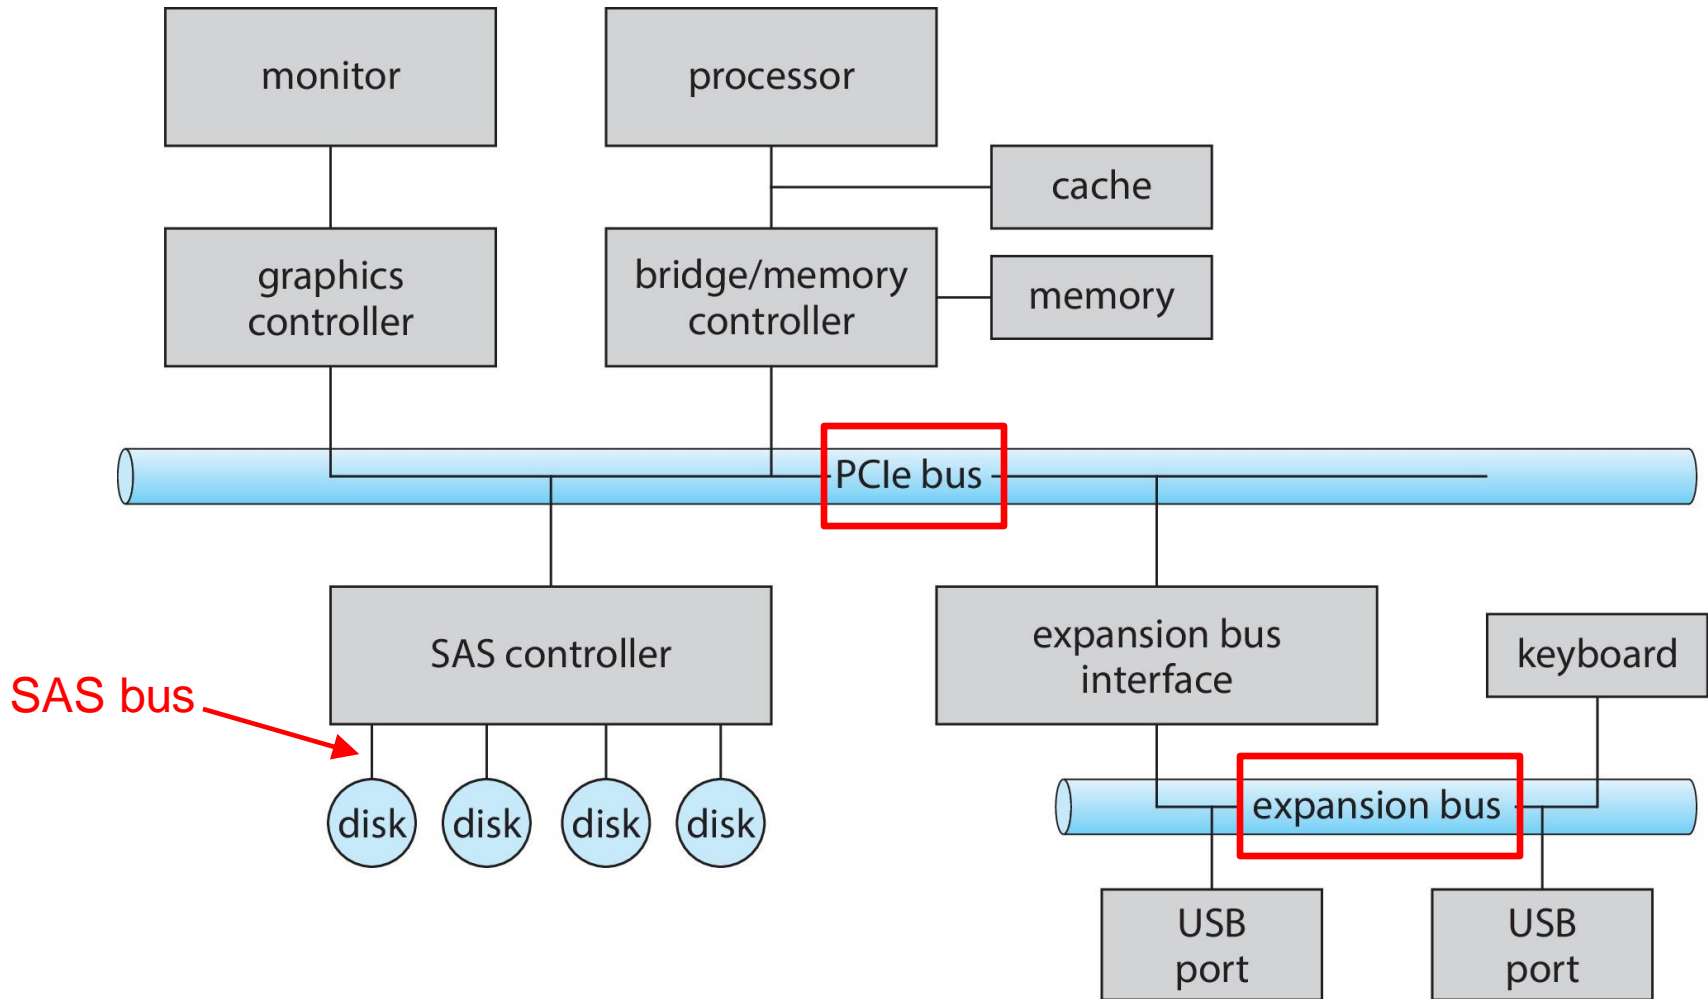


AR/VR Headset

I/O Hardware

- Incredible variety of I/O devices, including three general categories:
 - **Storage**: disks, tapes
 - **Transmission**: network connections, Bluetooth
 - **Human-interface**: screen, keyboard, mouse, audio in and out
- Common concepts – signals from I/O devices interface with computer
 - **Port** – connection point for devices
 - **Bus** - daisy chain or shared direct access
 - ▶ PCI bus common in PCs and servers, PCI Express (PCIe)
 - ▶ Expansion bus connects relatively slow devices
 - ▶ Serial-attached SCSI (SAS) common disk interface
 - **Controller (host adapter)** – electronics that operate port, bus, device
 - ▶ Sometimes integrated, sometimes separate circuit board (host adapter)
 - ▶ Contains processor, microcode, private memory, bus controller, etc.
 - ▶ **Fibre channel (FC)** is complex controller, usually separate circuit board (**host-bus adapter, HBA**) plugging into bus

A Typical PC Bus Structure



Host-Controller Communication

- I/O instructions control devices
- The controller has one or more registers for data and control signals.
 - E.g., data-in register, data-out register, status register, control register
 - Typically 1-4 bytes, or FIFO buffer
- Two ways of communication between the **host processor** and **controllers**:
 - **Option 1: Special I/O instructions**
 - ▶ Specify the transfer of a byte or a word to an I/O port address
 - **Option 2: Memory-mapped I/O**
 - ▶ Device-control registers are mapped into the processor address space
 - ▶ Especially for large address spaces (graphics)

Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)

Host-Controller Communication: Polling (轮询)

- The host writes output through a port, handshaking with the controller.
- For each byte of I/O,
 1. The host repeatedly reads `busy` bit from the `status` register until 0
 2. Host sets `write` bit in the `command` register, and write byte into `data-out` register
 3. Host sets `command-ready` bit
 4. Controller notices the `command-ready` bit, and sets `busy` bit
 5. Controller reads the `command` register, reads the `data-out` register, and does the I/O to the device
 6. Controller clears `busy` bit, `error` bit, `command-ready` bit when transfer done
- Step 1 is **busy-wait** cycle (or **polling**) to wait for I/O from device
 - Reasonable if device is fast, but inefficient if device is slow
 - If the wait may be long, CPU should switch to other tasks
 - ▶ How does the host know when the controller become idle?

Host-Controller Communication: Interrupts

- Polling can happen in 3 instruction cycles
 - 1) Read status, 2) logical-and to extract status bit, 3) branch if not zero
 - How to be more efficient if non-zero infrequently?
- **Interrupt**: The hardware controller notifies the CPU when the device is ready for service
 1. **Raised** by the device controller by asserting a signal on the interrupt request
 2. **Caught** by the CPU
 3. **Dispatched** to the interrupt handler
 4. **Cleared** by the interrupt handler by serving the device
- **Interrupt vector**: saves memory addresses of interrupt handlers
 - Two types of interrupt:
 - ▶ **Non-maskable**: Signal error conditions, page faults, and debug requests
 - ▶ **Maskable**: Used for device-generated interrupts
 - Interrupt chaining if more than one device at the same interrupt number
 - ▶ Each element points to the head of an interrupt handler list

Frequent Interrupts in OS

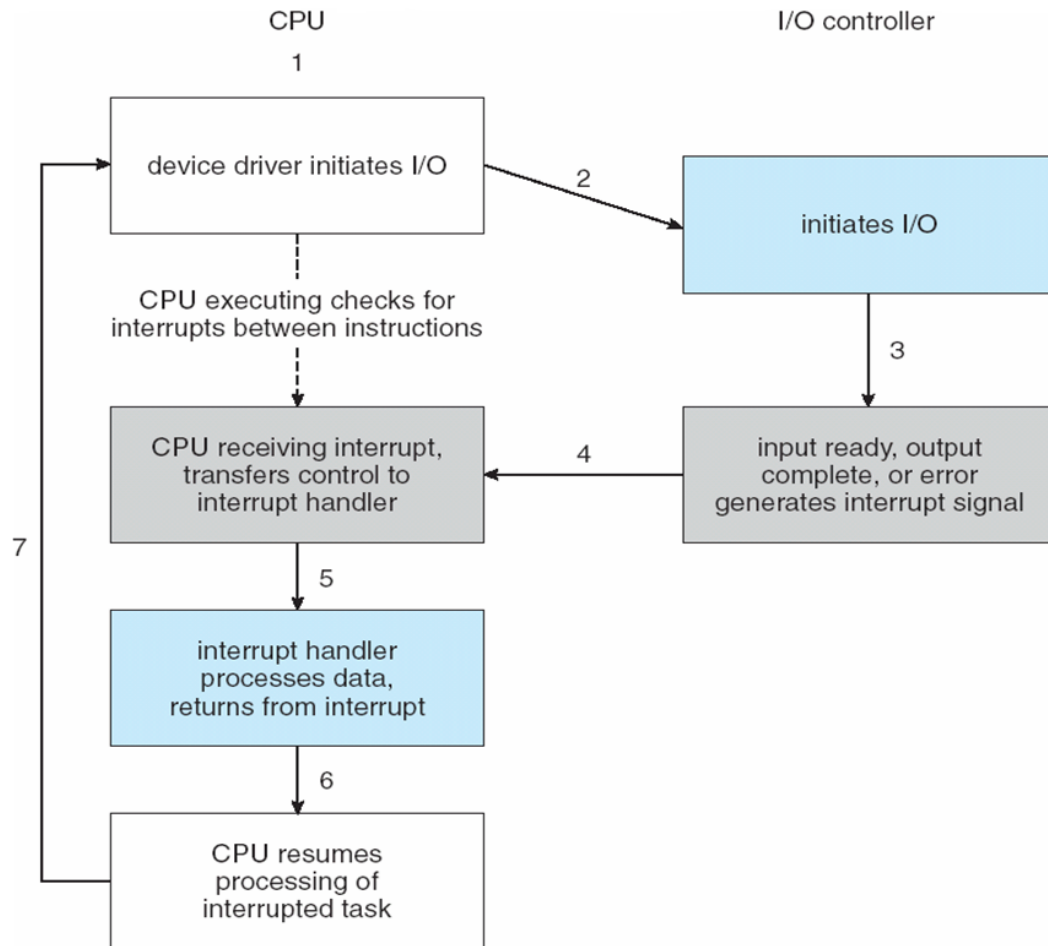
- **Stressing interrupt management** because
 - Single-user systems manage hundreds of interrupts per second
 - Servers can have hundreds of thousands of interrupts per second
- For example, a quiet macOS desktop generated 23,000 interrupts over 10 seconds

Fri Nov 25 13:55:59		0:00:10	
	SCHEDULER	INTERRUPTS	

total_samples	13	22998	
delays < 10 usecs	12	16243	
delays < 20 usecs	1	5312	
delays < 30 usecs	0	473	
delays < 40 usecs	0	590	
delays < 50 usecs	0	61	
delays < 60 usecs	0	317	
delays < 70 usecs	0	2	
delays < 80 usecs	0	0	
delays < 90 usecs	0	0	
delays < 100 usecs	0	0	
total < 100 usecs	13	22998	

Interrupt-Driven I/O Cycle

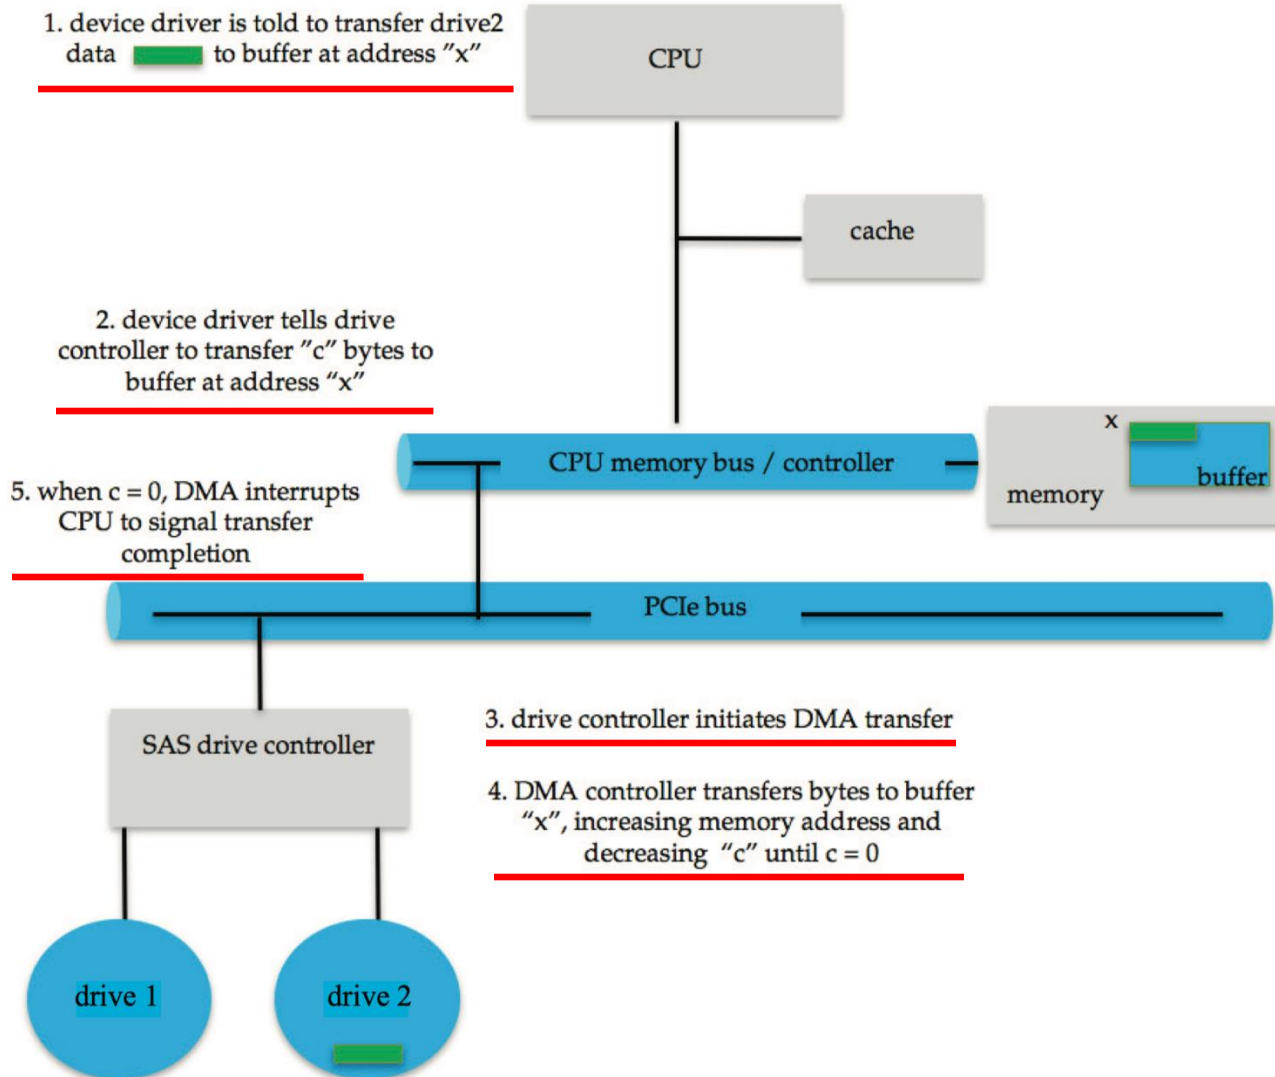
1. CPU senses **interrupt-request line** after executing every instruction
2. When CPU detects a controller has asserted a signal on the interrupt-request line, it performs a state save and jumps to the interrupt-handler
3. The **interrupt handler** determines the cause of the interrupt, performs the necessary processing
4. The interrupt handler performs a state restore, and returns the CPU to the execution state prior to the interrupt



Direct Memory Access (DMA)

- ❑ Used to avoid **programmed I/O** (one byte at a time) for host-controller communication with large data movement
 - ❑ It requires a **DMA** controller
- ❑ **Idea:** Bypasses CPU to transfer data directly between I/O device and memory
- ❑ To initiate a DMA transfer, OS writes DMA command block into the memory
 - ❑ The command block contains the source and destination addresses of the transfer, as well as the number of bytes to transfer
 - ❑ CPU writes the location of the command block to the DMA controller, then works on other tasks
 - ❑ **DMA controller operates the memory bus directly without CPU involvement**
 - ❑ When done, the DMA controller interrupts the CPU to signal completion
- ❑ Version that is aware of virtual addresses can be even more efficient - **DVMA**

Six Step Process to Perform DMA Transfer



I/O Hardware Concepts

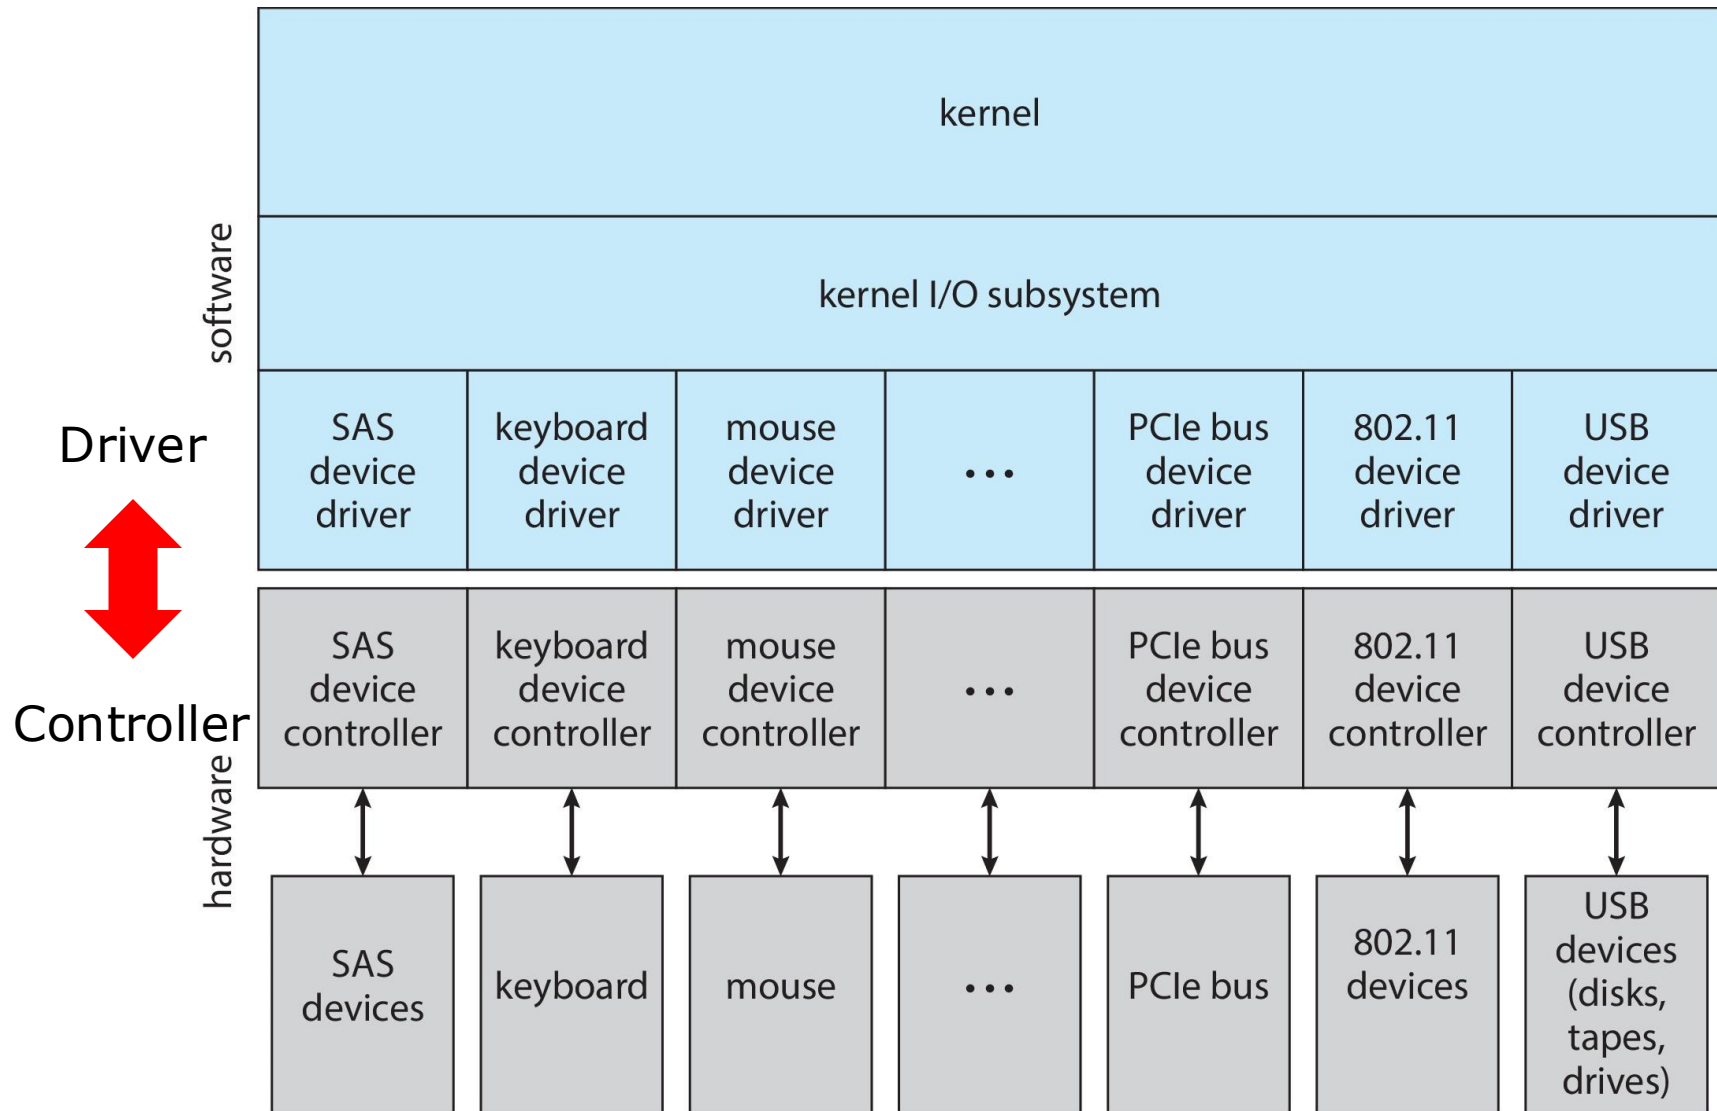
- The main concepts in I/O hardware:
 - A bus
 - A controller
 - An I/O port and its registers
- Handshaking between the **host** and a **device controller** through a polling loop or via interrupts
- Communication can be offloaded to a **DMA controller** for large data transfers

Application I/O Interface

Application I/O Interface

- ❑ I/O system calls encapsulate device behaviors in generic classes
 - ❑ Device-driver layer hides **I/O controllers differences** from the **kernel**
 - ❑ New devices using existing protocols need no extra work
- ❑ Each OS has own **I/O subsystem structures** and **device driver** frameworks
- ❑ Devices vary in many dimensions
 - ❑ Character-stream or block
 - ❑ Sequential or random-access
 - ❑ Synchronous or asynchronous (or both)
 - ❑ Sharable or dedicated
 - ❑ Speed of operation
 - ❑ read-write, read only, or write only

A Kernel I/O Structure



Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read–write	CD-ROM graphics controller disk

Characteristics of I/O Devices (Cont.)

- Differences of devices handled by device drivers
- Broadly I/O devices can be grouped by the OS into
 - Block I/O
 - Character I/O (Stream)
 - Memory-mapped file access
 - Network sockets
- For direct manipulation of I/O device specific characteristics, usually an escape / back door to transparently pass commands from application to device driver
 - Unix `ioctl()` [**I/O Control**] call to send arbitrary bits to a device control register and data to device data register
- UNIX and Linux use tuple of “**major**” and “**minor**” device numbers to identify type and instance of devices (here major 8 and instance number 0-4)

% `ls -l /dev/sda*`

```
brw-rw---- 1 root disk 8, 0 Mar 16 09:18 /dev/sda
brw-rw---- 1 root disk 8, 1 Mar 16 09:18 /dev/sda1
brw-rw---- 1 root disk 8, 2 Mar 16 09:18 /dev/sda2
brw-rw---- 1 root disk 8, 3 Mar 16 09:18 /dev/sda3
```

I/O Devices: (1) Block and Character Devices

- The **block-device** interface captures all the aspects necessary for accessing disk drives and other block-oriented devices.
- Block devices include **disk drives**
 - Commands include read, write, seek
 - Raw I/O, direct I/O, or file-system access
 - Memory-mapped file access can be on top of block-device drivers
 - ▶ File mapped to virtual memory and clusters brought via demand paging
- **Character devices** include keyboards, mouse, serial ports
 - Commands include **get()**, **put()**
 - Libraries layered on top allow line editing

I/O Devices: (2) Network Devices

- ❑ Performance differs from block and character to have their own interface
 - ❑ Facilitate distributed applications
- ❑ Linux, Unix, Windows , and many others include **socket** interface
 - ❑ Separates network protocol from network operation
 - ❑ Includes **select()** function
 - ▶ Returns which sockets have a packet waiting and which have room to accept new packet
- ❑ Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)


I/O Devices: (3) Clocks and Timers

- Provide current time, elapsed time, and timer to trigger operation
- Normal resolution about 1/60 second
 - Some systems provide higher-resolution timers
- **Programmable interval timer** used for timings, periodic interrupts
 - Can be set to wait a certain amount of time and then generate an interrupt
 - Disk I/O uses it to invoke periodic flushings of dirty cache buffers to disk
 - Network subsystem uses it to cancel operations that are proceeding too slowly because of network congestion or failures
- `ioctl()` (on UNIX) covers odd aspects of I/O such as clocks and timers

Blocking and Nonblocking I/O

- **Blocking I/O (阻塞)** - process suspended until I/O completed
 - Easy to use and understand
 - Insufficient for some needs
- **Nonblocking I/O (非阻塞)** - I/O call returns as much as available
 - User interface, data copy (buffered I/O)
 - Implemented via multi-threading
 - Returns quickly with the count of bytes read or written
 - Example: `select()` to find if data ready then `read()` or `write()` to transfer
- **Asynchronous** - process runs while I/O executes
 - Returns immediately, without waiting for the I/O to complete
 - I/O subsystem signals process when I/O completed
 - Not exposed to user applications but contained within the OS operation

Synchronous vs. Asynchronous I/O

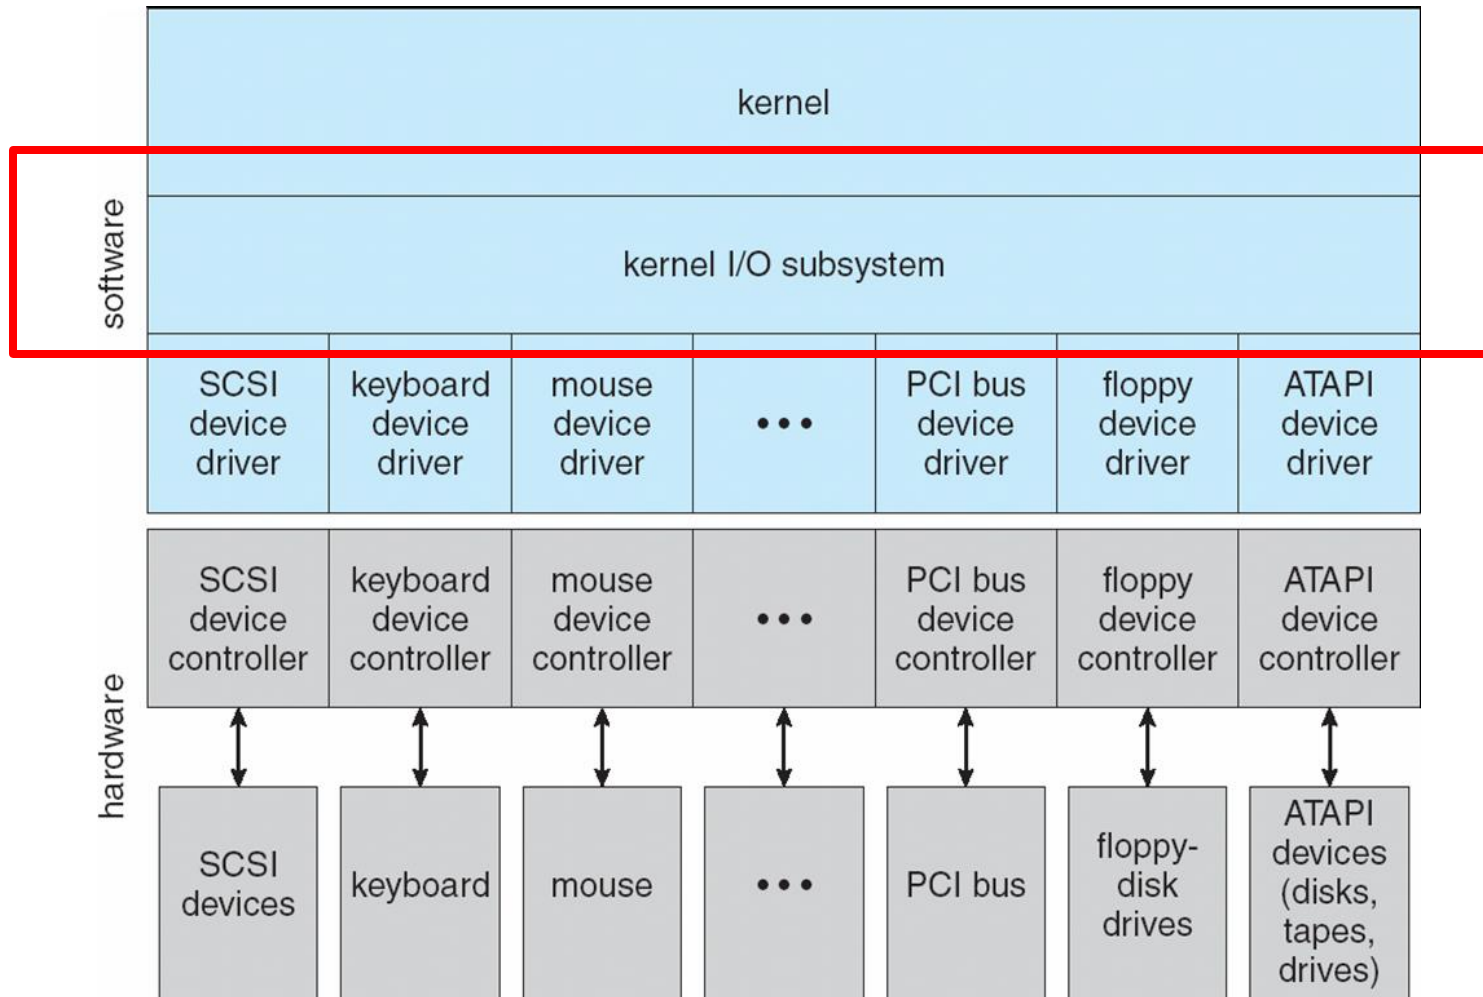


Vectored I/O

- **Vectored I/O (向量)** allows one system call to perform multiple I/O operations
- For example, Unix **readve()** accepts a vector of multiple buffers to read into or write from
- This scatter-gather method is better than multiple individual I/O calls
 - Decreases **context switching** and **system call** overhead
 - Some OS versions provide atomicity (**原子实现**)
 - ▶ No need to worry about multiple threads changing data as reads/writes occurring

Kernel I/O Subsystem

A Kernel I/O Structure



Kernel I/O Subsystem

- ❑ **I/O Scheduling** – Determine a good order to execute I/O requests
 - ❑ Some I/O request ordering via per-device queue
 - ❑ Some OSs try fairness
 - ❑ Some implement Quality of Service
- ❑ **Buffering** - Store data in memory while transferring between devices
 - ❑ To cope with device speed mismatch (e.g., network vs. disk)
 - ❑ To cope with device transfer size mismatch (e.g, different buffer sizes at sender and receiver)
 - ❑ To maintain “copy semantics”
- ❑ **Double buffering** – Two copies of the data
 - ❑ Decouples the producer of data from the consumer, thus relaxing timing requirements between them
 - ❑ Copy-on-write can be used for efficiency in some cases

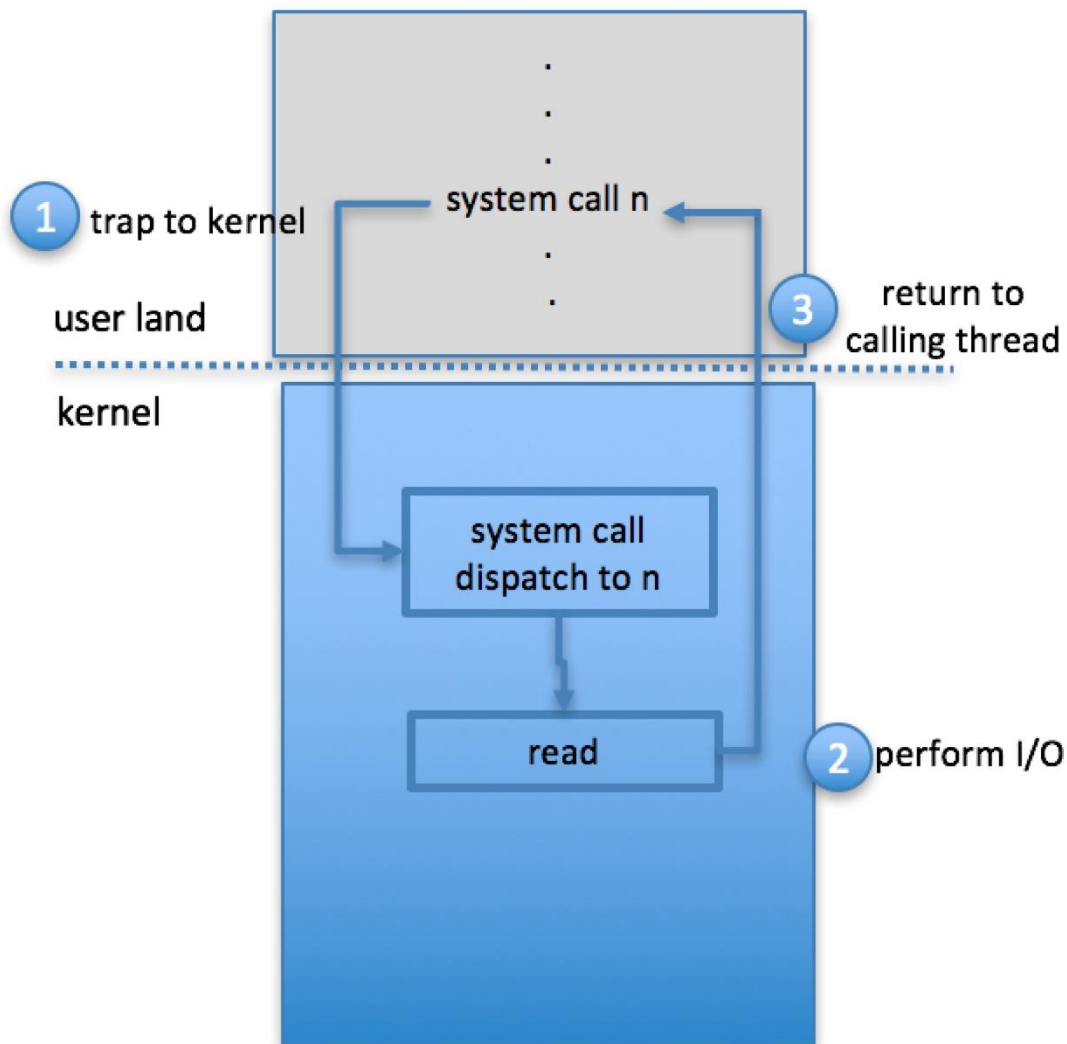
Kernel I/O Subsystem

- **Caching (缓存)** - faster memory region holding copy of data
 - Always just a copy
 - Key to performance
 - Sometimes combined with buffering
- **Spooling (假脱机)** - a buffer that holds output for a device
 - If device can serve only one request at a time
 - i.e., Printing
- **Device reservation** - provides exclusive access to a device
 - Enable a process to allocate an idle device and to deallocate that device when it is no longer needed
 - Watch out for deadlock

Error Handling and I/O Protection

- ❑ **Error handling:** OS can recover from disk read, device unavailable, transient write failures
 - ❑ Retry a read or write, for example
 - ❑ More advanced systems (Solaris FMA, AIX):
 - ▶ Track error frequencies
 - ▶ Stop using devices with increasing frequency of retry-able errors
- ❑ Most return an error number or code when I/O request fails
 - ❑ System error logs hold problem reports
- ❑ **I/O protection:** User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions
 - ❑ All I/O instructions defined to be privileged
 - ❑ I/O must be performed via system calls, but not by users
 - ▶ Memory-mapped and I/O port memory locations must be protected from user access

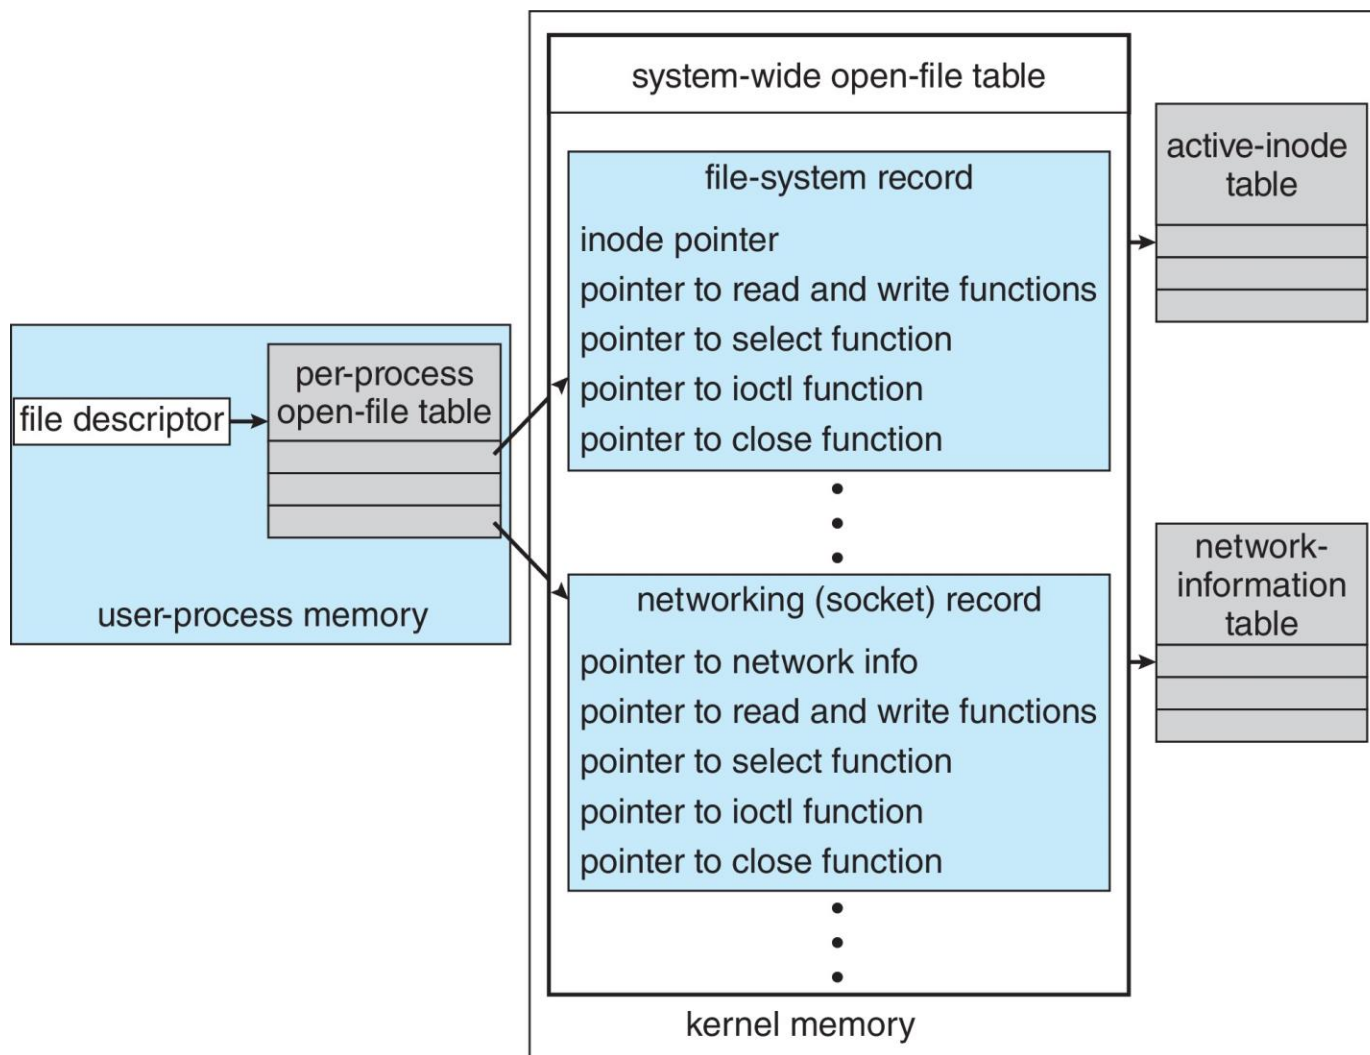
Use of a System Call to Perform I/O



Kernel Data Structures

- ❑ Kernel keeps state info for I/O components, including **open file tables**, **network connections**, **character device state**
- ❑ Many complex data structures to track buffers, memory allocation, “dirty” blocks
- ❑ Some use **object-oriented methods** and **message passing** to implement I/O
 - ❑ Windows uses message passing
 - ▶ Message with I/O information passed from user mode into kernel
 - ▶ Message modified as it flows through to device driver and back to process
 - ❑ Pros / cons?
 - ▶ Add overhead but simplify the structure and design of I/O system and adds flexibility

UNIX I/O Kernel Structure

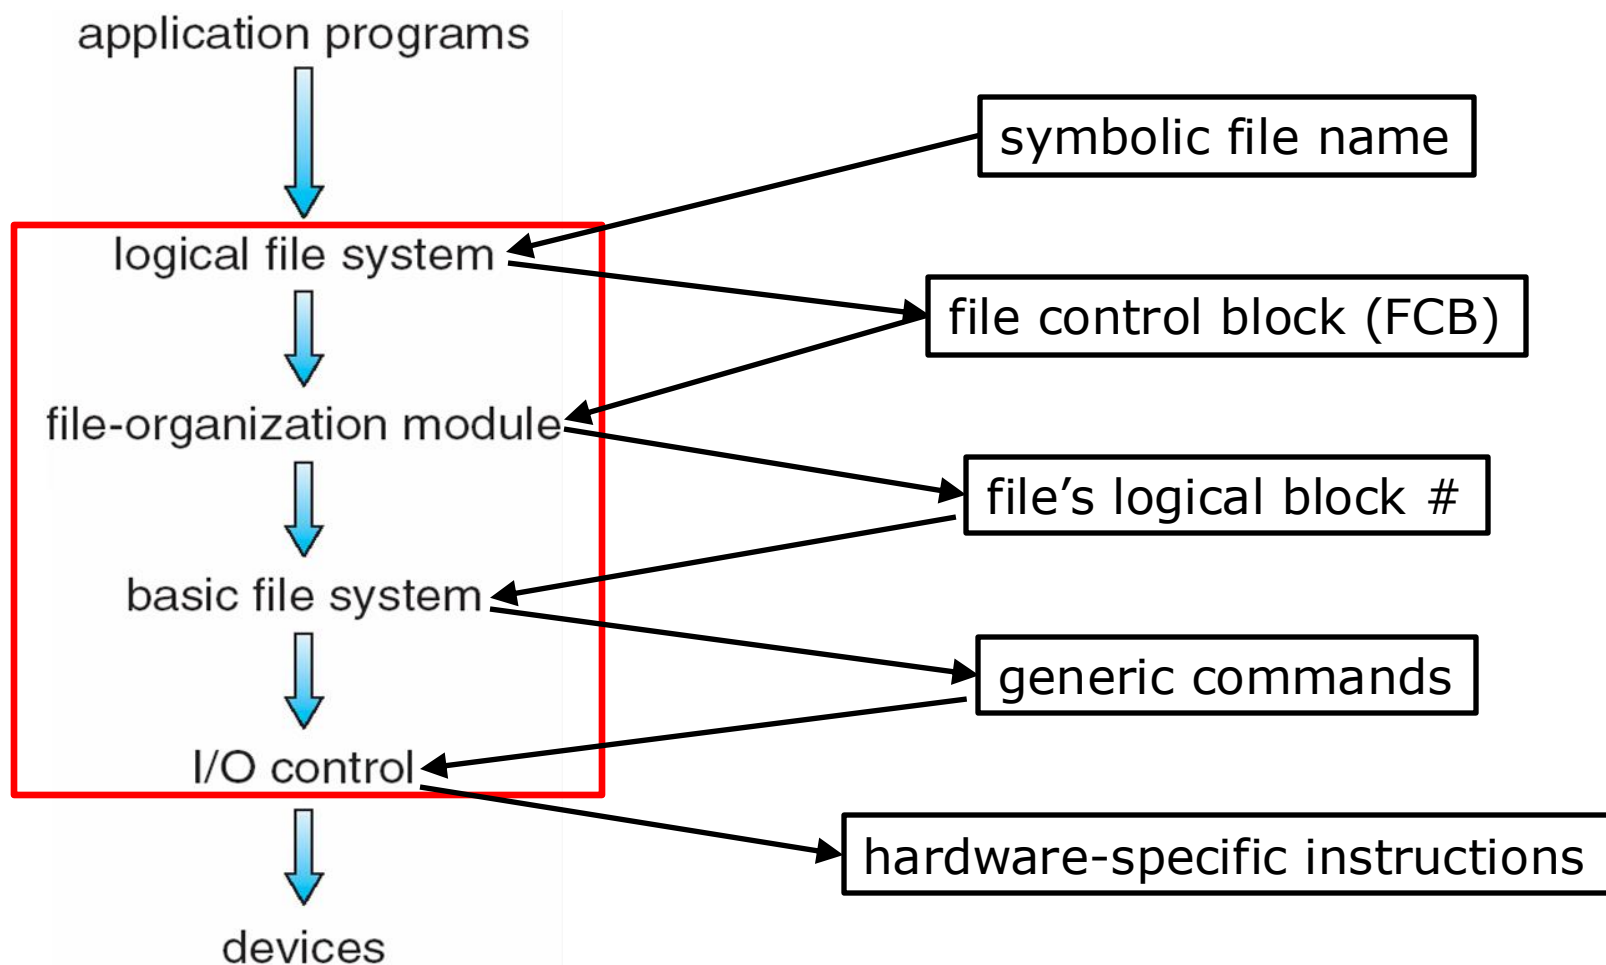


I/O Requests → Hardware Operations

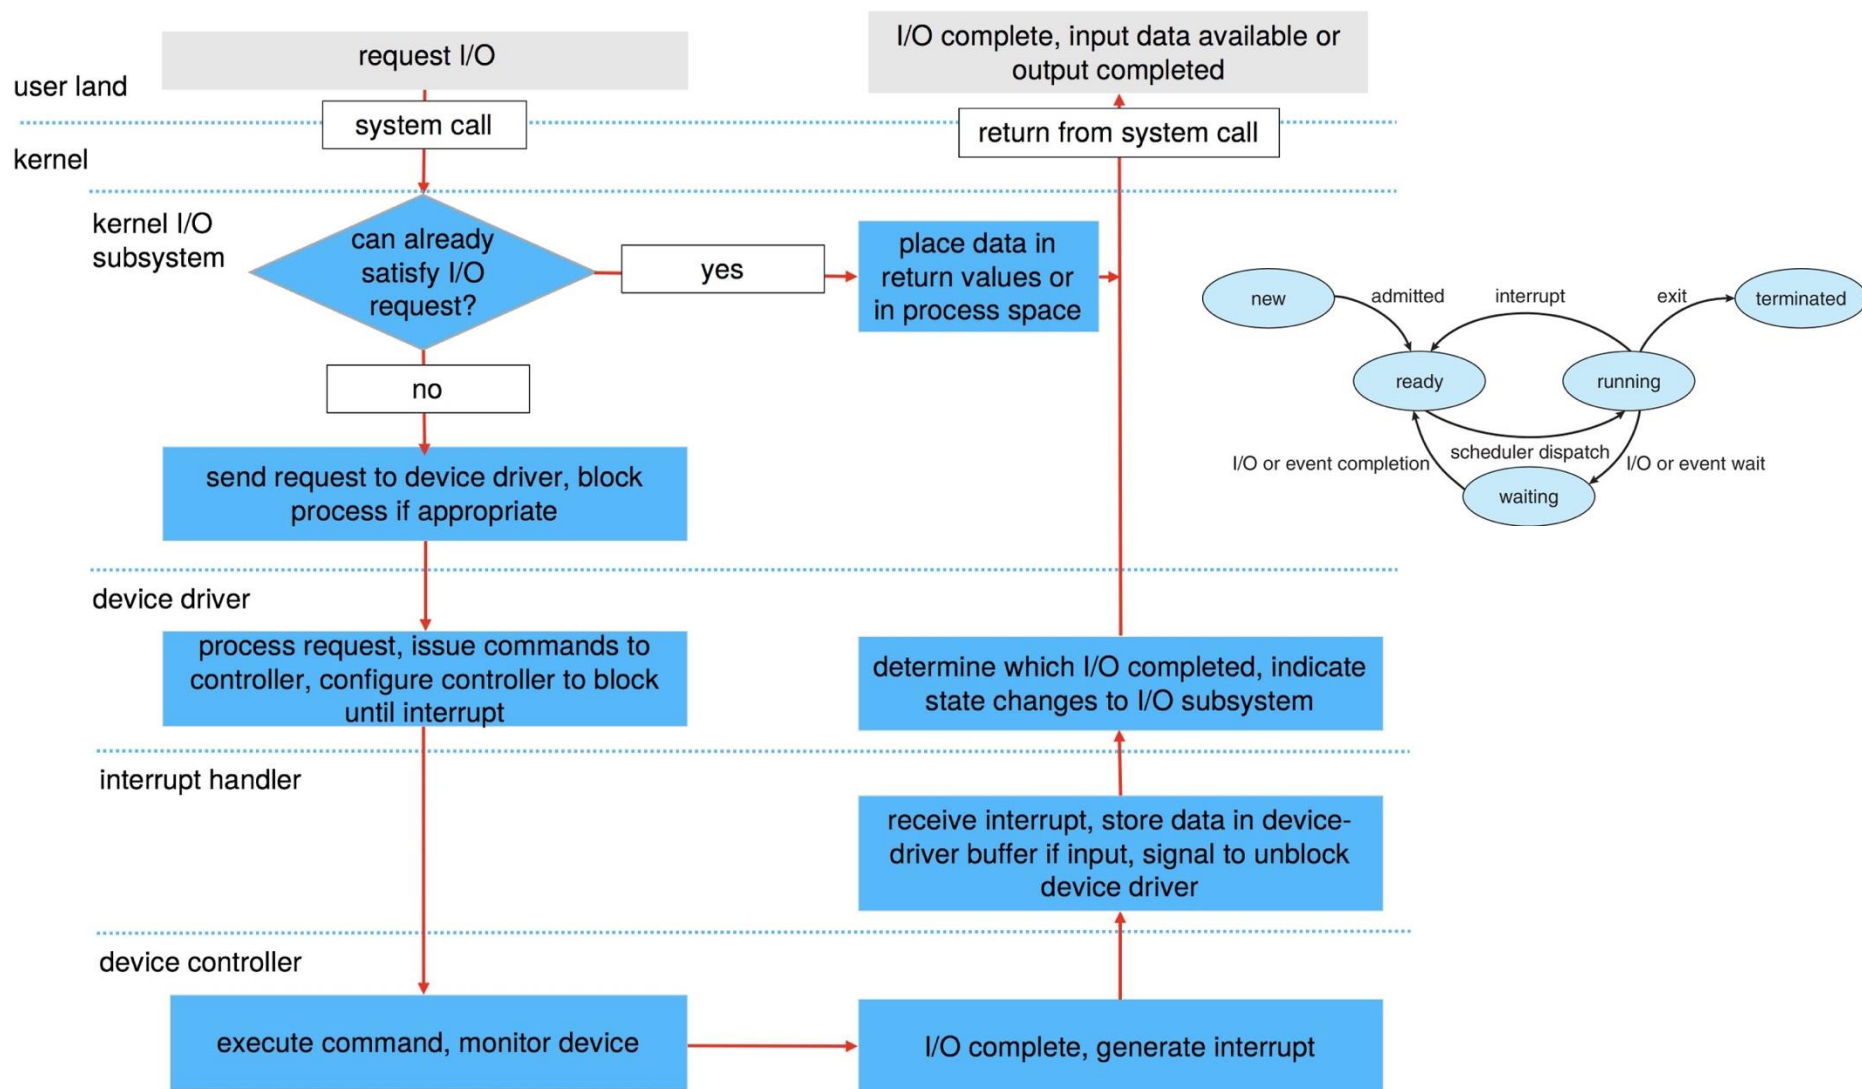
Transforming I/O Requests to Hardware Operations

- Consider **reading a file from disk** for a process:
 1. Determine the device holding the file
 2. **Translate name to device representation**
 3. Physically read data from disk into the buffer
 4. Make data available to the requesting process
 5. Return control to process

Recap: File System Layers and Operations



Life Cycle of An I/O Request



Summary

- The basic hardware elements involved in I/O are **ports, buses, device controllers**, and the **devices** themselves.
- Moving data between devices and main memory is performed by the CPU as **programmed I/O** or is **offloaded to a DMA controller**.
- The kernel module that controls a device is a **device driver**.
- The system call interface provided to applications is designed to handle several basic categories of hardware
 - Blocking vs. non-blocking vs. asynchronous calls
- Kernel's I/O subsystem provides numerous services: I/O scheduling, buffering, caching, spooling, device reservation, error handling, and translation.

Homework

- Reading
 - Chapter 12
- HW4 released today, due on **April 9, 23:59!**